



DeGirum AI Client for C++ Reference Guide

Contents

1	DeGirum AI Client for C++	1
1.1	Introduction	1
1.1.1	AI Server Package	1
1.1.2	AI Model Zoo	2
1.1.3	AI Client Package	2
1.2	AI Client Package Installation and Build	3
1.2.1	Installation Pre-Requisites	3
1.2.2	AI Client Package Installation	3
1.2.3	AI Client Package Build	4
1.3	AI Client C++ API Description	4
1.3.1	AI Server Connection and Addressing	4
1.3.2	Listing and Searching Models in Local Model Zoo	5
1.3.3	Working with Cloud Zoo Models	6
1.3.4	AIModel: Synchronous Sequential AI Inferencing API	7
1.3.5	Format of Inference Results	7
1.3.6	AIModelAsync: Asynchronous Pipelined AI Inferencing API	9
1.3.7	Model Parameters Management	11
1.3.8	Format of Data Frames for Image Processing Models	14
1.3.9	Error Handling	14
2	Hierarchical Index	17
2.1	Class Hierarchy	17
3	Class Index	19
3.1	Class List	19
4	File Index	21
4.1	File List	21
5	Namespace Documentation	23
5.1	DG Namespace Reference	23
5.1.1	Detailed Description	25
5.1.2	Enumeration Type Documentation	25
5.1.2.1	DetectionStatus	26
5.1.2.2	ServerType	26
5.1.3	Function Documentation	26
5.1.3.1	detectHostnameServers()	26
5.1.3.2	detectSubnetServers()	27
5.1.3.3	devCtrl()	27
5.1.3.4	errorCheck()	28
5.1.3.5	jsonGetMandatoryValue()	28
5.1.3.6	jsonGetOptionalValue()	28
5.1.3.7	jsonKeyExist()	29
5.1.3.8	jsonSetOptionalValue()	29
5.1.3.9	labelDictionary()	30
5.1.3.10	messagePrepare()	30

5.1.3.11	messagePrepareJson()	31
5.1.3.12	modelFind()	31
5.1.3.13	modelzooListGet()	31
5.1.3.14	modelZooManage()	32
5.1.3.15	serverPing()	32
5.1.3.16	shutdown()	33
5.1.3.17	systemInfo()	33
5.1.3.18	traceManage()	33
5.1.3.19	versionGet()	33
5.1.4	Variable Documentation	34
5.1.4.1	MODEL_PARAMS_CURRENT_VERSION	34
5.1.4.2	MODEL_PARAMS_MIN_COMPATIBLE_VERSION	34
5.1.4.3	SECT_TOP	34
6	Class Documentation	35
6.1	DG::AIModel Class Reference	35
6.1.1	Detailed Description	35
6.1.2	Constructor & Destructor Documentation	35
6.1.2.1	AIModel()	36
6.1.3	Member Function Documentation	36
6.1.3.1	predict()	36
6.2	DG::AIModelAsync Class Reference	37
6.2.1	Detailed Description	37
6.2.2	Member Typedef Documentation	37
6.2.2.1	callback_t	38
6.2.3	Constructor & Destructor Documentation	38
6.2.3.1	AIModelAsync()	38
6.2.3.2	~AIModelAsync()	39
6.2.4	Member Function Documentation	39
6.2.4.1	lastError()	39
6.2.4.2	predict()	39
6.2.4.3	setCallback()	39
6.2.4.4	waitCompletion()	40
6.3	DG::BasicTensor Class Reference	40
6.3.1	Detailed Description	42
6.3.2	Constructor & Destructor Documentation	43
6.3.2.1	BasicTensor() [1/3]	43
6.3.2.2	BasicTensor() [2/3]	43
6.3.2.3	BasicTensor() [3/3]	44
6.3.3	Member Function Documentation	44
6.3.3.1	adopt()	44
6.3.3.2	alloc() [1/3]	45
6.3.3.3	alloc() [2/3]	45
6.3.3.4	alloc() [3/3]	46
6.3.3.5	convert() [1/2]	46
6.3.3.6	convert() [2/2]	46
6.3.3.7	copy()	47
6.3.3.8	data() [1/2]	47
6.3.3.9	data() [2/2]	47
6.3.3.10	dequantize()	47
6.3.3.11	isEqualData() [1/2]	47
6.3.3.12	isEqualData() [2/2]	48
6.3.3.13	isEqualDataShape()	48
6.3.3.14	isEqualStruct()	48
6.3.3.15	linearSizeCalc()	48
6.3.3.16	quantize()	49
6.3.3.17	quantParamsSet()	49
6.3.3.18	reinterpretShapeForYolo()	49

6.3.3.19	reinterpretShapeScaled()	49
6.3.3.20	reshapeCombineDims()	50
6.3.3.21	reshapeTo() [1/2]	50
6.3.3.22	reshapeTo() [2/2]	50
6.3.3.23	shapeStringGet() [1/2]	50
6.3.3.24	shapeStringGet() [2/2]	50
6.4	DG::JsonHelper Class Reference	51
6.4.1	Detailed Description	51
6.4.2	Member Function Documentation	51
6.4.2.1	errorCheck()	51
6.4.2.2	jsonDeserialize()	52
6.4.2.3	jsonDeserializeStr()	52
6.4.2.4	jsonSerialize()	53
6.4.2.5	jsonSerializeStr()	53
6.4.2.6	parse()	53
6.4.2.7	parse_ignore_errors()	54
6.4.2.8	tensorDeserialize()	54
6.4.2.9	tensorSerialize()	54
6.5	DG::ModelInfo Struct Reference	55
6.5.1	Detailed Description	56
6.6	DG::ModelParams< Base, > Class Template Reference	56
6.6.1	Detailed Description	57
6.6.2	Constructor & Destructor Documentation	57
6.6.2.1	ModelParams() [1/4]	57
6.6.2.2	ModelParams() [2/4]	58
6.6.2.3	ModelParams() [3/4]	58
6.6.2.4	ModelParams() [4/4]	58
6.7	DG::ModelParamsReadAccess Class Reference	59
6.7.1	Detailed Description	60
6.7.2	Constructor & Destructor Documentation	60
6.7.2.1	ModelParamsReadAccess()	60
6.7.3	Member Function Documentation	60
6.7.3.1	fullJsonGet()	60
6.7.3.2	jsonGet()	61
6.7.3.3	modelShapeGet()	61
6.7.3.4	operator const json &()	61
6.7.3.5	operator std::string()	61
6.7.3.6	paramExist()	62
6.7.3.7	paramGet()	62
6.7.3.8	sectionGet()	63
6.7.3.9	sectionHashGet()	63
6.7.3.10	sectionSizeGet()	64
6.8	DG::ModelParamsSection Struct Reference	64
6.8.1	Detailed Description	64
6.9	DG::ModelParamsWriteAccess Class Reference	65
6.9.1	Detailed Description	66
6.9.2	Constructor & Destructor Documentation	66
6.9.2.1	ModelParamsWriteAccess()	66
6.9.3	Member Function Documentation	66
6.9.3.1	merge()	66
6.9.3.2	paramSet()	67
6.9.3.3	set_dirty()	67
6.10	DG::ModelQuery Struct Reference	67
6.10.1	Detailed Description	68
6.10.2	Member Enumeration Documentation	68
6.10.2.1	TriState	68
6.11	DG::BasicTensor::quant_params_t Struct Reference	69
6.11.1	Detailed Description	69

6.11.2	Constructor & Destructor Documentation	69
6.11.2.1	quant_params_t() [1/3]	69
6.11.2.2	quant_params_t() [2/3]	70
6.11.2.3	quant_params_t() [3/3]	70
6.11.3	Member Function Documentation	70
6.11.3.1	isEqualData()	70
6.11.3.2	isEqualStruct()	71
6.12	DG::BasicTensor::quant_params_t::scale_t Struct Reference	71
6.12.1	Detailed Description	71
6.13	DG::ServerAddress Struct Reference	71
6.13.1	Detailed Description	72
6.13.2	Constructor & Destructor Documentation	72
6.13.2.1	ServerAddress()	72
6.13.3	Member Function Documentation	73
6.13.3.1	fromHostname()	73
6.13.3.2	is_valid()	73
7	File Documentation	75
7.1	dg_client_structs.h File Reference	75
7.1.1	Detailed Description	77
7.2	dg_json_helpers.h File Reference	77
7.2.1	Detailed Description	78
7.3	dg_model_api.h File Reference	79
7.3.1	Detailed Description	80
7.4	dg_model_parameters.h File Reference	80
7.4.1	Detailed Description	82
7.5	dg_tensor_structs.h File Reference	82
7.5.1	Detailed Description	84
Index		85

Chapter 1

DeGirum AI Client for C++

1.1 Introduction

The DeGirum AI Client package referenced in this document is the part of DeGirum AI Software Suite.

Note: AI Client package consists of multiple programming language support packs. This document describes AI Client for C++. AI client packs for other languages are described in separate documents.

The DeGirum AI Software Suite provides means to perform fast and efficient AI model inferencing on locally installed AI hardware accelerator platforms.

The list of supported hardware accelerator platforms includes all DeGirum Orca AI hardware accelerators and Google Edge TPU accelerators. AI inferencing on local host CPU is also supported.

The DeGirum AI Software Suite consists of DeGirum AI Server package, AI Model Zoo, and DeGirum AI Client package.

1.1.1 AI Server Package

The AI Server package is installed on a computer system where AI hardware accelerators are connected to. The purpose of the AI Server is the following:

- It controls AI hardware accelerators
- It provides distributed and concurrent access of multiple AI clients to AI inferencing functionality of AI hardware accelerators over the TCP/IP network
- It maintains local AI model zoo with AI models to be used for AI inferencing
- It provides Web-based GUI to manage local model zoo

The following computer platforms are currently supported for the server installation:

- Intel x64 with Linux 64-bit OS
- ARM64 with Linux 64-bit OS

The support for following computer platforms is currently under development:

- Intel x64 with Windows 10 OS
- Intel x64 with MacOS
- ARM64 with MacOS

The AI Server for Linux is distributed as fully configured Docker container (refer to <https://hub.docker.com/r/degirum/aiserver>), which simplifies AI Server deployment to the variety of Linux platforms.

1.1.2 AI Model Zoo

The AI Model Zoo is the place where AI models reside.

The local AI Model Zoo of DeGirum AI Server is implemented as the directory on the AI Server local file system. It contains files of all locally deployed AI models. Only models in the local Model Zoo can be used for AI inferencing on particular AI server installation.

Currently, AI Server supports the following AI model types:

- N2X models: models in proprietary DeGirum N2X format compiled for DeGirum Orca AI accelerator hardware and CPU
- TFLite Edge TPU models: models in Google TFLite format compiled for Google Edge TPU AI accelerator hardware
- TFLite CPU models: models in Google TFLite format compiled for execution on CPU
- OpenVINO models

The deployment of the local Model Zoo is performed using PySDK server module (please refer to PySDK documentation available on <https://degirum.github.io>).

1.1.3 AI Client Package

The AI Client package is installed on a computer system, where client-side ML software runs. It can be the same computer system, where the AI Server is deployed, or it can be any other computer system, which has TCP/IP network access to the AI Server computer system.

AI Client Package includes the following components:

- AI Client library
- AI Client examples
- Build scripts

AI Client library provides APIs to control and use the AI Server. It includes the following functionality:

- List all AI models with their attributes available for AI inferencing in particular AI Server local Model Zoo
- Find AI model in that local Model Zoo matching particular model attributes
- Open and configure streaming connection to particular AI Server for inferencing using particular AI model from local Model Zoo
- Stream data frames to the that streaming connection to perform AI inferencing on these data frames

AI Client library provides integration into the C++ programming language. This document describes AI Client for C++. AI client packs for other languages are described in separate documents.

All AI Client Package software is distributed in the source code and requires compilation.

The following computer platforms are currently supported for the client installation:

- Intel x64 with Linux 64-bit OS
- Intel x64 with Windows 10 OS
- ARM64 with Linux 64-bit OS
- Intel x64 with MacOS
- ARM64 with MacOS

1.2 AI Client Package Installation and Build

1.2.1 Installation Pre-Requisites

For Linux OS the following software packages need to be installed in order to build AI Client package software:

- Git
- CMake
- GCC compiler
- Thread building blocks library

To install all these component use the following commands:

```
sudo apt install git-all cmake g++ libtbb-dev
```

For Windows OS the following software packages need to be installed in order to build AI Client package software:

- Git
- CMake
- Visual Studio 2019 (Community Edition would suffice)

1.2.2 AI Client Package Installation

The AI Client library is distributed in source code and is available on GitHub at the following link:

```
https://github.com/DeGirum/CppSDK
```

You clone the AI Client library repo by executing command:

```
git clone https://github.com/DeGirum/CppSDK
```

The repo directory contains the AI Client library collaterals. The subdirectory structure is as follows:

Directory	Contents
/client	AI Client library code
/doc	documentation
/examples	sample applications
/extern	third-party libraries
/images	sample images to be used for AI inference
/inc	AI Client library C++ include files

1.2.3 AI Client Package Build

The CMakeLists.txt file is provided for the CMake build. It builds AI Client C++ library, and C++ examples.

To perform the build under Linux OS, execute the following commands:

```
cd client
mkdir build
cd build
cmake ..
cmake --build .
```

For Windows OS and Visual Studio, CMake will create the Visual Studio project files and AIClientLib.sln solution.

For Linux OS, CMake will create make files (unless otherwise specified with the -G flag).

After building, the library and example executables are copied into /bin subdirectory.

1.3 AI Client C++ API Description

DeGirum AI Client library provides the following functionality:

- [versionGet\(\)](#) function to query the version of the AI Client library
- [modelzooListGet\(\)](#) function to list all AI models with their attributes available for AI inferencing in particular AI Server local Model Zoo
- [systemInfo\(\)](#) function to obtain AI server system information including information about connected AI devices
- [modelFind\(\)](#) function to find AI model in that local Model Zoo matching particular model attributes
- [labelDictionary\(\)](#) function to query the model class label dictionary
- [errorCheck\(\)](#) function to check given AI server JSON response for run-time errors
- [DG::AIModel](#) class to connect to AI server and to perform AI model inference in a simple non-pipelined sequential manner (see [AIModel: Synchronous Sequential AI Inferencing API](#))
- [DG::AIModelAsync](#) class to connect to AI server and to perform AI model inference in efficient pipelined asynchronous manner using mechanism of callbacks (see [AIModelAsync: Asynchronous Pipelined AI Inferencing API](#))
- the family of ModelParams classes to handle model parameters (see [Model Parameters Management](#))
- functions and methods for error handling (see [Error Handling](#))

1.3.1 AI Server Connection and Addressing

The client software communicates with AI server through TCP/IP network using DeGirum proprietary communication protocol based on ASIO library. In order to establish a connection to AI server, the client side must know the host name or IP address of the AI server, and TCP port, which AI server listens.

The server address is passed to AI inference classes as a string in the following format:

If AI server domain name is known, the address string is specified in the form

```
domain_name:port
```

If only IP address of AI server is known, the address string is specified in the form

```
xxx.xxx.xxx.xxx:port
```

The port number together with ':' separator can be omitted, in this case default port 8778 is used. When AI server is started without specifying the port, the same default port 8778 is used for the server side. In this case it is safe to omit the port. Otherwise you should match the port on the client side to the port used on the server side.

1.3.2 Listing and Searching Models in Local Model Zoo

The AI server can perform AI inferences only on models located either in the local model zoo of that AI server or in the cloud zoo you have access to. This section describes how to deal with local zoo models. The section [Working with Cloud Zoo Models](#) describes how to deal with cloud zoo models.

To initiate an inference of a model from the local model zoo on AI server, the client must provide to the server the *model name*, which uniquely identifies the desired model in the local model zoo. There are two ways for the client code to obtain the model name.

1. By obtaining the full list of models available in the model zoo using [modelzooListGet\(\)](#) function
2. By searching the model zoo for a model satisfying certain search attributes using [modelFind\(\)](#) function.

Either way you obtain the instance of [DG::ModelInfo](#) structure, which contains all necessary model identification information, including:

- [DG::ModelInfo::name](#) - model string name
- [DG::ModelInfo::extended_params](#) - extended model parameters (all parameter as listed in model JSON configuration file)

Then you pass model name [DG::ModelInfo::name](#) to constructors of classes used for AI inference: [DG::AIModel](#) or [DG::AIModelAsync](#).

The [modelzooListGet\(\)](#) function returns an array of [ModelInfo](#) structures listing all models found in the local model zoo. It is then your responsibility to iterate over all elements of this array to find the model you want to use.

The [modelFind\(\)](#) function provides simpler means to find the desired model. It accepts the model query: the instance of [DG::ModelQuery](#) structure, which contains some model attributes you want to look for. The [modelFind\(\)](#) function first obtains the full list of models from the AI server, and then iterates over this list to find a model matching all attributes specified in the model query. It returns the first model matching all attributes.

You may specify the following attributes:

- Any part of the model name. Only models which have the specified substring in their model name would be considered.
- Device type to use for inference. Only models which are designed for specified device type would be considered.
- Runtime agent to use for inference. Only models which are designed for specified runtime agent would be considered.
- Model should be quantized flag. Only quantized models would be considered.
- Model should be pruned/sparse flag. Only sparse/pruned models would be considered.

Only the model name substring attribute is mandatory, all other attributes are optional and can be omitted if not needed.

You specify the device type by providing the Device ID string. The following device types are supported:

Device ID	Device
"ORCA"	DeGirum Orca AI accelerator
"EDGETPU"	Google Edge TPU AI accelerator
"MYRIAD"	Intel Myriad AI accelerator
"CPU"	AI server host computer

The *runtime agent* is the software module, which controls the device used for AI inference. You specify the runtime agent type by providing the Agent ID string. The following runtime agent types are supported:

Agent ID	Agent
"N2X"	DeGirum N2X runtime agent. Used for all Orca devices and N2X model inference on CPU.
"TFLITE"	Google TFLite runtime agent. Used for Edge TPU devices and direct TFLite model inference on CPU.
"OPENVINO"	Intel OpenVINO runtime agent. Used for Myriad devices and direct OpenVINO model inference on CPU.

Note: the runtime agent attribute is rarely required because in most use cases the AI hardware uniquely defines the runtime agent to use, so it is recommended to avoid setting this attribute.

Examples:

Find a model which has 'MobileNet_v2' substring in the model name:

```
auto model_id = DG::modelFind( server, { "mobilenet_v2" } );
```

Find a model which has 'mobilenet_v2' substring in the model name, is designed to run on DeGirum Orca AI accelerator hardware, is quantized, and is sparse/pruned:

```
auto model_id = DG::modelFind( server, { "MobileNet_v2",  
    "ORCA", "", DG::ModelQuery::Yes, DG::ModelQuery::Yes } );
```

1.3.3 Working with Cloud Zoo Models

Starting from release 0.5.0, the AI server supports inference of models from DeGirum-hosted cloud model zoos. In order to use models from a cloud zoo you need to register at DeGirum AI Hub, <https://hub.degirum.com>, and generate API access token. Please refer either to PySDK documentation or General Information section, both available on <https://hub.degirum.com>, for more details.

To do the inference of a model from some cloud zoo you need to perform the following steps:

1. Obtain your **API access token**. You can generate it on DeGirum AI Hub site <https://hub.degirum.com> under Management | My Tokens* main menu item.
2. Figure out the **extended model name**. It has the following format: "<organization>/<zoo>/<model>". The "<organization>" field is the name of the zoo owner organization, the "<zoo>" field is the name of the cloud zoo, and the "<model>" field is the name of the model in that cloud zoo. You can obtain the "<organization>/<zoo>" prefix (so-called model zoo path) on DeGirum AI Hub site <https://hub.degirum.com> under Management | Models main menu item. Just select the model zoo you want to access: the model zoo page opens. Then click the Copy button near the model zoo name to copy the model zoo path into the clipboard. All models available in the cloud zoo are also listed on the model zoo page. Click on particular model: the model page opens. The model name will be displayed at the top of the page. Now concatenate the model zoo path with the model name separated by the slash / to obtain the extended model name.
3. Create `DG::ModelParamsWriter` object and assign `CloudToken` parameter.
4. Create `DG::AIModel` or `DG::AIModelAsync` object passing the AI server hostname, extended model name, and `DG::ModelParamsWriter` object.

For example: `const std::string hostname = ""; // specify AI server hostname`
`const std::string ext_model_name = ""`
`// specify extended model name`
`const std::string token = ""; // specify your API access token string`

```

DG::ModelParamsWriter mparams; mparams.CloudToken_set( token );
DG::AIModel model( hostname, ext_model_name, mparams );

```

AI server will automatically download the model from the cloud model zoo using provided credentials and store this model in the local model cache to be reused for later inferences.

Each time you create new `DG::AIModel` or `DG::AIModelAsync` object and perform the first inference, the checksum of the model in the local cache is compared with the checksum of the model in the cloud zoo, and if they mismatch, the model from the cloud zoo is downloaded into the local cache. This mechanism guarantees that you always perform the inference of the newest model.

1.3.4 AIModel: Synchronous Sequential AI Inferencing API

The `DG::AIModel` class is used to perform AI model inference on AI server in simple non-pipelined sequential manner.

Note: for more efficient (but somewhat more complex to use) asynchronous pipelined inference use `DG::AIModelAsync` class.

On construction this class performs connection to AI server, selection of AI model, and setting model run-time parameters.

You specify the AI server address as a string as described in the section [AI Server Connection and Addressing](#).

You specify the model name as described in the section [Listing and Searching Models in Local Model Zoo](#).

If you want to modify some of the model parameters (see more on that in the section [Model Parameters Management](#)), you define the instance of `ModelParamsWriteAccess`-derived class, change model parameters you want to adjust by invoking setter methods, and pass this instance as one of `DG::AIModel` constructor arguments.

For example, you want to change the input frame image type to JPEG. You can pass the following expression to `DG::AIModel` constructor:

```

DG::ModelParamsWriter().InputImgFmt_set( "JPEG" )

```

Once constructed, it can be used to perform sequential AI inference by invoking `DG::AIModel::predict()` method. `Predict` method accepts the input frame data in various formats (see section [Format of Data Frames for Image Processing Models](#) for more details) and returns the inference results, which are described in the section [Format of Inference Results](#). This method is blocking method, i.e. it returns the execution only when the inference of an input frame is complete. Working with blocking methods simplifies the programming, but loses some run-time efficiency because (unlike for `DG::AIModelAsync` class) some actions like sending data to AI server, frame pre-processing, result post-processing, and receiving results from AI server are not pipelined between frames.

1.3.5 Format of Inference Results

The `predict()` method of `DG::AIModel` class returns JSON array with inference results. The internal structure of this JSON array greatly depends on the AI model type: it is different for classification models, detection models, and so on. But in general, that JSON array is an array of identical JSON objects, and each JSON object contains inference results for one detected entity.

The following tables describe JSON array format for some model types.

For classification model types each classified object is represented by the following JSON object:

JSON Field Name	Description	Data Type
"label"	Classified object class label	string
"category_id"	Classified object class index	integer
"score"	Probability of classified object	floating point

For detection model types each detected object is represented by the following JSON object:

JSON Field Name	Description	Data Type
"label"	Detected object class label	string
"category_id"	Detected object class index	integer
"score"	Probability of detected object	floating point
"bbox"	Detected object rectangle	four-element JSON array [xmin, ymin, xmax, ymax]
"landmarks"	Landmark points (optional)	JSON array of two-element arrays [[x1, y1], [x2, y2], ...]

For pose detection types each detected object is represented by the following JSON object:

JSON Field Name	Description	Data Type
"score"	Probability of detected pose	floating point
"landmarks"	Pose description landmarks	array of landmark JSON objects

The following is the structure of pose description landmark JSON object:

JSON Field Name	Description	Data Type
"label"	Classified object class label	string
"category_id"	Classified object class index	integer
"score"	Probability of classified landmark	floating point
"landmark"	landmark point	two-element JSON array [x, y]
"connect"	list of adjacent landmarks indices	JSON array of integers

For hand detection types each detected object is represented by the following JSON object:

JSON Field Name	Description	Data Type
"score"	Probability of detected hand	floating point
"handedness"	Probability of right hand	floating point
"landmarks"	hand description landmarks	array of landmark JSON objects

The following is the structure of hand description landmark JSON object:

JSON Field Name	Description	Data Type
"label"	Classified object class label	string
"category_id"	Classified object class index	integer
"landmark"	landmark point	three-element JSON array [x, y, z]
"world_landmark"	metric world landmark point	three-element JSON array [x, y, z]
"connect"	list of adjacent landmarks indices	JSON array of integers

When you specify post-processing type as "None" (see section [Model Parameters Management](#)) then the inference

result JSON array contains network output tensors. Each element of that array corresponds to one tensor. Each tensor is represented by the following JSON object:

JSON Field Name	Description	Data Type
"id"	Tensor numeric ID as specified in network	integer
"name"	Tensor name as specified in network	string
"shape"	Tensor shape: sizes of each dimension	integer array
"quantization"	Tensor quantization parameters	object
"type"	Tensor element type	string
"data"	Tensor linear data buffer contents	binary field

The following tensor data types are supported:

Type String	Type Description
DG_FLT	32-bit floating point
DG_UINT8	8-bit unsigned integer

The following is the structure of quantization JSON object:

JSON Field Name	Description	Data Type
"axis"	Quantization axis or -1 for global quantization	integer
"scale"	Quantization scale array	floating point array
"zero"	Quantization zero offset	integer array

For convenient handling of tensors the [DG::BasicTensor](#) class is provided. You can convert JSON array containing a tensor into [DG::BasicTensor](#) object by using [DG::JsonHelper::tensorDeserialize\(\)](#) static method:

```
for( const auto &json_tensor: inference_result )
{
    DG::BasicTensor basic_tensor( DG::JsonHelper::tensorDeserialize( json_tensor ) );
    auto shape = basic_tensor.shape();           // get tensor shape
    auto element_type = basic_tensor.dataTypeGet(); // get tensor element type
    if( element_type == DG_UINT8 )                // get pointer to tensor linear buffer
        const uint8_t *data = basic_tensor.data< uint8_t >();
}
```

There is a special case of JSON array returned by AI server in case when runtime error happens during inference. In this case the runtime error info object is returned. It is described in details in the [Error Handling](#) section.

1.3.6 AIModelAsync: Asynchronous Pipelined AI Inferencing API

The [DG::AIModelAsync](#) class is used to perform AI model inference on AI server in efficient pipelined asynchronous manner using the mechanism of callbacks.

Note: for simple (but less efficient) synchronous non-pipelined inference use [DG::AIModel](#) class.

On construction this class performs connection to AI server, selection of AI model, setting model run-time parameters, and installation of the client callback function. The client callback function is used to pass the inference results from the AI server to the client code.

You specify the AI server address as a string as described in the section [AI Server Connection and Addressing](#).

You specify the model name as described in the section [Listing and Searching Models in Local Model Zoo](#).

You should define the callback function somewhere in your code in the following form:

```
void MyAiCallback( const json &result, const std::string &info ){}
```

Then you pass your callback name (in this example it is `MyAiCallback`) as one of `DG::AIModelAsync` constructor arguments. The callback function has two arguments: the inference result and the frame info string. The inference result is JSON array, described in the section [Format of Inference Results](#). The frame info string is the string passed to `predict()` method along with the frame data. More on that below.

If you want to modify some of the model parameters (see more on that in the section [Model Parameters Management](#)), you define the instance of `DG::ModelParamsWriteAccess` derived class, change model parameters you want to adjust by invoking setter methods, and pass this instance as one of `DG::AIModelAsync` constructor arguments.

For example, you want to change the input frame image type to JPEG. You can pass the following expression to `DG::AIModelAsync` constructor:

```
DG::ModelParamsWriter().InputImgFmt_set( "JPEG" )
```

Once constructed, it can be used to perform asynchronous AI inference by invoking `DG::AIModelAsync::predict()` method. The `predict` method accept the input frame data in various formats (see section [Format of Data Frames for Image Processing](#) for more details) and initiates the inference on the AI server. The method is a non-blocking method, i.e. it returns execution immediately after posting the frame data to the AI server. This allows calling this method in a loop without waiting for the inference results, achieving the maximum AI server utilization.

To simplify matching inference results and corresponding frames in the client callback, additional argument can be passed to `predict()` method. The `frame_info` argument is optional frame information string. It will be passed to the client callback along with the frame result. You can pass arbitrary information as the frame info, for example frame data filenames, or frame numbers, or any other frame identification information, which you need to know in the scope of client callback function in order to interpret inference results in respect to the input frames.

Once the inference of a frame is complete, and the inference result is received from the AI server, the client callback is invoked to dispatch the inference result. Such result handling via callback mechanism is performed in a thread, separate from the main execution thread. It means that the client callback function is called in asynchronous manner, thus the name of the class.

During construction you may also specify the depth of the internal frame queue. The internal frame queue works the following way. If `predict()` method is invoked too often and the number of non-processed ("outstanding") frames exceeds the specified queue depth, the consecutive call to `predict()` method will be blocked until the number of outstanding frames in the queue becomes smaller than the queue depth thus allowing to post one more frame.

The class destructor waits until all outstanding results are received and then closes the connection to AI server. You may use `DG::AIModelAsync::waitCompletion()` to explicitly wait for completion of all outstanding inferences. It is a blocking call: it returns only when all outstanding frames are processed by AI server and all results are dispatched via client callback.

To monitor the number of outstanding inferences in real-time you use `DG::AIModelAsync::outstandingResultsCountGet()` method.

To check for runtime errors which happened during inferences you use `DG::AIModelAsync::lastError()` method. If ever during consecutive calls to `predict()` method the AI server reported a run-time error, then this method will return the error message string, otherwise it returns an empty string. Alternatively, you may check for runtime errors directly in the callback function by invoking `errorCheck()` function on inference result argument.

Note: in case of server runtime error, all frames posted after the error was detected, will not be processed.

1.3.7 Model Parameters Management

Each AI model in the model zoo comes with the model parameters file in JSON format. This JSON file defines all model parameters such as model name, model input dimensions, pre-processing parameters, post-processing parameters, etc. Some of the model parameters are fixed and cannot be altered without breaking the model behavior, and some of the model parameters can be changed to adjust the model behavior. Later parameters are called *runtime parameters* and can be adjusted when constructing instances of model inference classes [DG::AIModel](#) and [DG::AIModelAsync](#).

To work with model parameters the following classes are provided:

- [DG::ModelParamsReadAccess](#)
- [DG::ModelParamsWriteAccess](#)
- [DG::ModelParams](#)
- [DG::ModelParamsWriter](#)

[DG::ModelParamsReadAccess](#) class provides programmatic type-safe read access to model parameters defined in JSON model configuration. It keeps non-owning const reference to underlying JSON array. For each model parameter it provides getter method, which name matches the parameter name as it appears in JSON array.

[DG::ModelParamsWriteAccess](#) inherits [DG::ModelParamsReadAccess](#) and provides programmatic type-safe read and write access to model parameters defined in JSON model configuration. It keeps non-owning non-const reference to underlying JSON array. For each model parameter it provides both getter and setter methods. Getter method name matches the parameter name as it appears in JSON array. Setter method name is constructed from the parameter name by appending `_set` suffix. For example, `InputImgFmt_set()` Setter methods return reference to self, so they can be called in a daisy-chain manner, for example:

```
params.InputImgFmt_set( "RAW" ).InputImgRawDataType_set( "DG_UINT8" );
```

[DG::ModelParams](#) class provides programmatic type-safe access to model parameters defined in JSON model configuration while *owning* that JSON array. Access (read or write) is defined by the template parameter, which is used as the base class. You can use any [DG::ModelParamsReadAccess](#)-derived base class: it can be [DG::ModelParamsReadAccess](#) or [DG::ModelParamsWriteAccess](#).

And finally [DG::ModelParamsWriter](#) is [DG::ModelParams](#) template instantiation with write access. Instances of this class you typically use to pass adjusted runtime model parameters to constructors of inference classes.

The following table summarizes all runtime parameters which are accessible to the client code:

Parameter Name	Type	Access	Description	Possible Values	Default Value
CloudToken	string	R/W	Cloud API access token	Valid cloud API access token	""
EagerBatchSize	size_t	R/W	The size of the batch to be used by device scheduler when inferencing this model. The batch is the number of consecutive frames before this model is switched to another model during batch predict.	1..80	8
InputImgFmt	string	R/W	Input frame format: JPEG or raw binary	"JPEG", "RAW"	"JPEG"

Parameter Name	Type	Access	Description	Possible Values	Default Value
InputRawData↔ Type	string	R/W	Data type for raw binary frame format↔ : byte of 32-bit float	"DG_UINT8", "DG_↔FLT"	"DG_UINT8"
InputImg↔ Rotation	int	R/W	Input image rotation angle in degrees, clockwise	0, 90, 180, 270	0
InputColor↔ Space	string	R/W	Input image colorspace (sequence of colors in C dimension)	"RGB", "BGR"	"RGB"
InputResize↔ Method	string	R/O	Interpolation algorithm used for image resizing during model training	"nearest", "bilinear", "area", "bicubic", "lanczos"	"bilinear"
InputPadMethod	string	R/O	How input image is padded when resized during model training	"stretch", "letterbox", "crop-first", "crop-last"	"letterbox"
InputCrop↔ Percentage	double	R/O	How much input image was cropped during model training	0..1	1
ImageBackend	string	R/O	Graphical package used for image processing during model training	"pil", "opencv", "auto"	"auto"
MaxDetections	int	R/W	Maximum number of objects to report for detection models	> 0	20
Max↔ Detections↔ PerClass	int	R/W	Maximum number of objects to report for each class for detection models	> 0	100
MaxClasses↔ PerDetection	int	R/W	Maximum number of classes to report for detection models	> 0	30
MeasureTime	bool	R/W	Enable inference time measurements	true/false	false
Output↔ Postprocess↔ Type	string	R/W	Post-processor type	see below	"None"
OutputConf↔ Threshold	double	R/W	Confidence threshold	0..1	0.1
OutputNMS↔ Threshold	double	R/W	Threshold for non-max suppression	0..1	0.6
OutputTopK	size↔ _t	R/W	Number of classes with biggest scores to report for classification models	0,1...	0
UseRegularNMS	bool	R/W	Use regular (per-class) NMS algorithm as opposed to global (class-ignoring) N↔MS algorithm for detection models	true/false	true

If you do not specify a runtime parameter, then the value for this parameter will be taken from the model JS↔

ON configuration file from the local model zoo. If the model JSON configuration file also does not specify this parameter, then the default value will be used.

The following table lists currently available post-processor types.

Post-Processor Name	Applicability	Description
"None"	All models	Pass-through post-processor. Network output tensors are sent without any post-processing
"Classification"	Classification models	Classification post-processor for all classification models
"Segmentation"	Segmentation models	Semantic segmentation post-processor for all segmentation models
"TFLiteDetection"	TFLite detection models	Object detection post-processor for TFLite detection models
"DetectionYolo"	Yolov5-based models	Object detection post-processor for DeGirum Yolov5-based models
"Detection"	Other detection models	Object detection post-processor for all other DeGirum general object detection models
"DetectionYoloPlates"	Yolov5-based license plate detection models	DeGirum license plate detection model post-processor
"DetectionYoloV8Plates"	Yolov8-based license plate detection models	DeGirum license plate detection model post-processor
"DetectionYoloV8OBB"	Yolov8-based oriented bounding box (OBB) detection models	OBB object detection post-processor for DeGirum Yolov8-based OBB models
"FaceDetection"	Face detection models	Face detection post-processor for all DeGirum face detection models
"PoseDetection"	Pose detection models	Person pose detection post-processor for all DeGirum pose detection models
"HandDetection"	Hand detection models	Hand landmark detection post-processor for all DeGirum hand detection models
"Dequantize"	For quantized outputs	Dequantize output tensors and provide float tensors as the output

Note: not all post-processors are compatible with all models. Special care should be taken when specifying post-processor type other than "None".

Notes of Input* parameters.

DG::ModelParamsReadAccess::InputImgFmt parameter is used to tell pre-processor what input frame format to expect. The value of this parameter should match the frame data type you pass to predict() methods. You may change it to match the format of your input frames.

DG::ModelParamsReadAccess::InputRawDataType parameter is used to tell pre-processor what data element type to expect for raw data frames (InputImgFmt is "RAW"). Again, the value of this parameter should match the frame data element type you pass to predict() methods. You may change it to match the format of your input frames.

DG::ModelParamsReadAccess::InputImgRotation parameter is used to ask pre-processor to perform input image rotation by certain angle. It is convenient when your input frames are coming from your camera which is mounted on a side or upside down.

DG::ModelParamsReadAccess::InputColorSpace defines the colorspace (pixel order) which the AI model expects (was trained to). In case when the input image format is "RAW", you should provide raw image frame with correct pixel order corresponding to the input colorspace parameter value. In case when the input image format is "JPEG", the pre-processor will perform the conversion to the proper pixel order when decoding the JPEG image. Usually you do not change this parameter, otherwise the model performance will degrade.

1.3.8 Format of Data Frames for Image Processing Models

The AI Client package supports two formats of the input data frames for image processing models: JPEG images and binary bitmaps, where binary format is used by default and typically set in JSON model configuration files for all DeGirum models.

In any case the HxW dimension of the input data frame must match the input dimension of the AI model used for the inference. The input dimensions of the AI model are stored in the `DG::ModelInfo` data structure members `DG::ModelInfo::extended_params.InputW()` and `DG::ModelInfo::extended_params::InputH()`, returned by either `modelzooListGet()` or `modelFind()` functions (see section [Listing and Searching Models in Local Model Zoo](#)).

The frame data in the binary bitmap format should be arranged as HWC three-dimensional dense array, meaning that color dimension 'C' is the fastest dimension, while the vertical dimension 'H' is the slowest. The color order is dependent on how the model was originally trained. The pixel data type can be either 8-bit unsigned integer number or 32-bit floating point number.

You specify input frame type when you construct instances of model inference classes `DG::AIModel` and `DG::AIModelAsync`, passing instance of `ModelParams` class as the constructor parameter (see more details in the section [Model Parameters Management](#)).

To specify JPEG frame type you need to set `InputImgFmt` parameter to "JPEG".

To specify binary frame type you need to set both `InputImgFmt` and `InputImgRawDataType` parameters. For example, the byte binary frame type you set this way:

```
std::string server = "192.168.0.1"; // some server address
auto model_id = DG::modelFind( server, { "MobileNet_v2" } );
DG::AIModel model( server, model_id,
    DG::ModelParamsWriter.InputImgFmt_set( "RAW" ).InputImgRawDataType_set( "DG_UINT8" ) );
```

1.3.9 Error Handling

There are two mechanisms of error handling employed in AI Client package: using standard library exceptions and using specially formatted JSON inference results.

All runtime errors *except errors happened during inference*, are handled by throwing standard library exceptions - objects of `std::exception` type. To handle such errors the recommended practice is to place all the code which interacts with the AI Client library into one try/catch block, and then catch exceptions of `std::exception` type:

```
try
{
    ...
}
catch( std::exception &e )
{
    std::cout << e.what(); // handle exception
}
```

Errors which happened during the inference are handled special way. Instead of throwing an exception, the special JSON runtime error info object is returned by `DG::AIModel::predict()` method in case you use synchronous API, and that JSON object is passed as an argument into your callback in case you use asynchronous API.

The JSON runtime error info object has the following structure:

JSON Field Name	Description
"success"	Boolean field containing "false" in case of error
"msg"	Error message string

If you want to check that the inference result actually defines runtime error, you use [errorCheck\(\)](#) function. If the AI server JSON response contains the run-time error, then this function will return the error message string, otherwise it returns an empty string. The following code demonstrates the typical use case for synchronous API:

```
json result;
model.predict( frame, result );
auto possible_error = DG::errorCheck( result );
if( !possible_error.empty() )
    throw std::runtime_error( possible_error );
```

The following code demonstrates the typical use case for asynchronous API:

```
void MyAiCallback( const json &result, const std::string &info )
{
    auto possible_error = DG::errorCheck( result );
    if( possible_error.empty() )
    {
        ... accept inference result
    }
    else // do not throw exception, since it will be ignored by AI Client thread
}
```

Note: do not throw any exceptions in your callback, since the AI Client callback thread, which handles user callback invocation, suppresses all exceptions thrown in callbacks to not to terminate the callback thread unexpectedly.

Another approach for the inference error handling for asynchronous API would be to check for inference errors at the very end of the inference sequence. You can use [DG::AIModelAsync::lastError\(\)](#) method: if ever during consecutive calls to `predict()` method the inference errors are detected, then this method will return the error message string, otherwise it returns an empty string. To guarantee that all outstanding frames are processed before the error check, you can use [DG::AIModelAsync::waitCompletion\(\)](#) method:

```
model.waitCompletion();
auto possible_error = model.lastError();
if( !possible_error.empty() )
    throw std::runtime_error( possible_error );
```

Note: in case of inference error in asynchronous client, all frames posted after the error was detected, will not be processed.

Chapter 2

Hierarchical Index

2.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

DG::AIModel	35
DG::AIModelAsync	37
DG::BasicTensor	40
DG::JsonHelper	51
DG::ModelInfo	55
DG::ModelParamsReadAccess	59
DG::ModelParams< Base, >	56
DG::ModelParamsWriteAccess	65
DG::ModelParams< ModelParamsWriteAccess >	56
DG::ModelParamsSection	64
DG::ModelQuery	67
DG::BasicTensor::quant_params_t	69
DG::BasicTensor::quant_params_t::scale_t	71
DG::ServerAddress	71

Chapter 3

Class Index

3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

DG::AIModel	
AIModel	is DeGirum AI client API class for simple non-pipelined sequential inference 35
DG::AIModelAsync	
AIModelAsync	is DeGirum AI client API class for efficient pipelined asynchronous inference 37
DG::BasicTensor	 40
DG::JsonHelper	
JsonHelper	helper class: contains assorted static methods to operate with JSON configurations 51
DG::ModelInfo	
ModelInfo	is the model identification structure. It keeps AI model key attributes 55
DG::ModelParams< Base, >	
ModelParams	is model parameters collection with user-defined access rights 56
DG::ModelParamsReadAccess	
ModelParamsReadAccess	is read-only accessor to model parameters 59
DG::ModelParamsSection	
ModelParamsSection	model parameters section descriptor 64
DG::ModelParamsWriteAccess	
ModelParamsWriteAccess	is read/write accessor to model parameters 65
DG::ModelQuery	
ModelQuery	model query structure used to search models on AI server which match a set of provided model attributes 67
DG::BasicTensor::quant_params_t	
quant_params_t	Quantization parameters (how to convert from integer/quantized data back to floating point) 69
DG::BasicTensor::quant_params_t::scale_t	 71
DG::ServerAddress	 71

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

dg_client_structs.h	
DG client API data types	75
dg_json_helpers.h	
DG Core JSON helper classes and functions	77
dg_model_api.h	
DG Client API for model inference	79
dg_model_parameters.h	
DG centralized handling of JSON model parameters	80
dg_tensor_structs.h	
DG tensor container classes	82

Chapter 5

Namespace Documentation

5.1 DG Namespace Reference

[DG](#) is DeGirum namespace which groups all DeGirum software declarations.

Classes

- class [AIModel](#)
[AIModel](#) is DeGirum AI client API class for simple non-pipelined sequential inference.
- class [AIModelAsync](#)
[AIModelAsync](#) is DeGirum AI client API class for efficient pipelined asynchronous inference.
- class [BasicTensor](#)
- class [JsonHelper](#)
[JsonHelper](#) is JSON helper class: contains assorted static methods to operate with JSON configurations.
- struct [ModelInfo](#)
[ModelInfo](#) is the model identification structure. It keeps AI model key attributes.
- class [ModelParams](#)
[ModelParams](#) is model parameters collection with user-defined access rights.
- class [ModelParamsReadAccess](#)
[ModelParamsReadAccess](#) is read-only accessor to model parameters.
- struct [ModelParamsSection](#)
[ModelParamsSection](#) is model parameters section descriptor.
- class [ModelParamsWriteAccess](#)
[ModelParamsWriteAccess](#) is read/write accessor to model parameters.
- struct [ModelQuery](#)
[ModelQuery](#) is model query structure used to search models on AI server which match a set of provided model attributes.
- struct [ServerAddress](#)

Typedefs

- using [BasicTensorVector](#) = std::vector< [BasicTensor](#) >
Collection of tensors.
- using [json](#) = nlohmann::json
JSON library.
- typedef struct [DG::ModelInfo](#) [ModelInfo](#)
[ModelInfo](#) is the model identification structure. It keeps AI model key attributes.
- using [ModelParamsWriter](#) = [ModelParams](#)< [ModelParamsWriteAccess](#) >
[ModelParamsWriter](#) is [ModelParams](#) template instantiation with write access.
- using [vec3d_int](#) = std::vector< std::vector< std::vector< int > > >
3D vector of integers
- using [vec_dbl](#) = std::vector< double >
vector of doubles
- using [vec_flt](#) = std::vector< float >
vector of floats
- using [vec_int](#) = std::vector< int >
vector of integers
- using [vec_str](#) = std::vector< std::string >
vector of strings
- using [vec_sz](#) = std::vector< size_t >
vector of size_t

Enumerations

- enum [DetectionStatus](#) { [DetectionStatus::OK](#), [DetectionStatus::ProtocolMismatch](#), [DetectionStatus::HostAlive](#) }
Enumerator showing server status as detected by the below function.
- enum [ServerType](#) { [ServerType::Unknown](#), [ServerType::ASIO](#), [ServerType::HTTP](#) }
Server protocol type.

Functions

- std::vector< std::tuple< std::string, [DG::DetectionStatus](#) > > [detectHostnameServers](#) (const std::string &prefix, const int range_start, const int range_end, const int numeral_width=3)
- std::vector< std::tuple< std::string, [DetectionStatus](#) > > [detectSubnetServers](#) (const std::string &root_ip, const std::string &subnet_mask)
- [json devCtrl](#) (const std::string &server, const [json](#) &req)
- std::string [errorCheck](#) (const [json](#) &json_response)
- template<typename T >
T [jsonGetMandatoryValue](#) (const [json](#) &json_params, const std::string §ion, int index, const std::string &key)
- template<typename T >
T [jsonGetOptionalValue](#) (const [json](#) &json_params, const std::string §ion, int index, const std::string &key, const T &default_value)
- static bool [jsonKeyExist](#) (const [json](#) &json_params, const std::string §ion, int index, const std::string &key)
- template<typename T >
void [jsonSetOptionalValue](#) ([json](#) &json_params, const std::string §ion, int index, const std::string &key, const T &value)
- [json labelDictionary](#) (const std::string &server, const std::string &model_name)
- static std::string [messagePrepare](#) (const [json](#) &in)

- static `json messagePrepareJson` (const `json` &in)
- `ModelInfo modelFind` (const std::string &server, const `ModelQuery` &query)
- void `modelzooListGet` (const std::string &server, std::vector< `ModelInfo` > &modelzoo_list)
- `json modelZooManage` (const std::string &server, const `json` &req)
- std::ostream & `operator<<` (std::ostream &os, const `BasicTensor` &t)
Print tensor to stream.
- bool `serverPing` (const std::string &server)
- void `shutdown` (const std::string &server)
- `json systemInfo` (const std::string &server)
- `json traceManage` (const std::string &server, const `json` &req)
- std::string `versionGet` ()

Variables

- const int `CURRENT_PROTOCOL_VERSION` = 4
current client-server protocol version
- const int `DEFAULT_CONNECTION_TIMEOUT_MS` = 10000
Default connection timeout, ms.
- const int `DEFAULT_FRAME_QUEUE_DEPTH` = 8
Default frame queue depth.
- const int `DEFAULT_INFERENCE_TIMEOUT_MS` = 180000
Default inference timeout, ms.
- const int `DEFAULT_PORT` = 8778
Default TCP port of AI server.
- const int `MIN_COMPATIBLE_PROTOCOL_VERSION` = 4
minimum compatible client-server protocol version
- const int `MODEL_PARAMS_CURRENT_VERSION` = 12
- const int `MODEL_PARAMS_MIN_COMPATIBLE_VERSION` = 1
- constexpr const char * `PROTOCOL_VERSION_TAG` = "VERSION"
client-server protocol version tag
- constexpr `ModelParamsSection SECT_DEVICE` = { "DEVICE", true }
device parameters section
- constexpr `ModelParamsSection SECT_INTERNAL` = { "INTERNAL", true }
internal parameters section
- constexpr `ModelParamsSection SECT_MODEL_PARAMETERS` = { "MODEL_PARAMETERS", true }
model parameters section
- constexpr `ModelParamsSection SECT_POST_PROCESS` = { "POST_PROCESS", true }
post-processing parameters section
- constexpr `ModelParamsSection SECT_PRE_PROCESS` = { "PRE_PROCESS", false }
pre-processing parameters section
- constexpr `ModelParamsSection SECT_TOP` = { "", true }
JSON file top-level sections.

5.1.1 Detailed Description

`DG` is DeGirum namespace which groups all DeGirum software declarations.

5.1.2 Enumeration Type Documentation

5.1.2.1 DetectionStatus

```
enum DG::DetectionStatus [strong]
```

Enumerator showing server status as detected by the below function.

Enumerator

OK	Server fully operational.
ProtocolMismatch	Server is up, but is outdated.
HostAlive	Hardware is alive, but no server found. Unimplemented!

5.1.2.2 ServerType

```
enum DG::ServerType [strong]
```

Server protocol type.

Enumerator

Unknown	Not set.
ASIO	proprietary TCP socket server protocol
HTTP	HTTP server protocol.

5.1.3 Function Documentation

5.1.3.1 detectHostnameServers()

```
std::vector< std::tuple< std::string, DG::DetectionStatus > > DG::detectHostnameServers (
    const std::string & prefix,
    const int range_start,
    const int range_end,
    const int numeral_width = 3 )
```

Detect all ORCA servers with hostnames generated from a prefix and a range

Parameters

in	<i>prefix</i>	is a string prefix, formatted as "nnnn[:port]". This will be used to generate two sets of hostnames, one with the pattern "nnnn#", and one "nnnn###". For example, the prefix "farm", range_start 1, range_end 2 and numeral_width 3 will scan the hosts "farm1", "farm001", "farm2", "farm002". Setting numeral_width to 0 disables pattern generation, and will just scan the prefix. The port part is optional. If included, all servers will be scanned on the port. If omitted, the default port 8778 will be used.
----	---------------	--

Parameters

in	<i>range_start</i>	is the lowest numeral to be applied to the pattern
in	<i>range_end</i>	is the highest numeral to be applied to the pattern, inclusive
in	<i>numeral_width</i>	is the width of the numeral in the padded set. This is optional, set to 3 by default. Set to 0 to disable pattern generation.

Returns

Vector of strings containing hostnames of ORCA servers on the network.

5.1.3.2 detectSubnetServers()

```
std::vector< std::tuple< std::string, DetectionStatus > > DG::detectSubnetServers (
    const std::string & root_ip,
    const std::string & subnet_mask )
```

Detect all ORCA servers on a given subnet

Parameters

in	<i>root_ip</i>	is an IP address in the network, formatted as "xxx.xxx.xxx.xxx[:port]". The port part is optional. If included, all servers will be scanned on the port. If omitted, the default port 8778 will be used.
in	<i>subnet_mask</i>	specifies the subnet with the associated IP, formatted as "xxx.xxx.xxx.xxx", where "xxx" is either 255 for fixed bits or 0 for variable bits.

Returns

Vector of pairs containing IPs of ORCA servers on the network and their statuses

5.1.3.3 devCtrl()

```
json DG::devCtrl (
    const std::string & server,
    const json & req )
```

Device control facility

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	<i>req</i>	- management request

Returns

results of execution of requests

5.1.3.4 errorCheck()

```
std::string DG::errorCheck (
    const json & json_response )
```

Check given AI server JSON response for run-time errors. If AI server reported run-time error in the given JSON response, then this function extracts and returns the error message string from the provided JSON array, otherwise it returns an empty string.

5.1.3.5 jsonGetMandatoryValue()

```
template<typename T >
T DG::jsonGetMandatoryValue (
    const json & json_params,
    const std::string & section,
    int index,
    const std::string & key )
```

Get value from JSON array. Raise error, if not there

Parameters

in	<i>json_params</i>	- JSON array
in	<i>section</i>	- top section name, which contains array of records; if empty, key is taken from topmost level
in	<i>index</i>	- index in array of records
in	<i>key</i>	- key of key-value pair

Returns

parameter value

5.1.3.6 jsonGetOptionalValue()

```
template<typename T >
T DG::jsonGetOptionalValue (
    const json & json_params,
    const std::string & section,
    int index,
    const std::string & key,
    const T & default_value )
```

Get value from JSON array, which may not be present there

Parameters

in	<i>json_params</i>	- JSON array
in	<i>section</i>	- top section name, which contains array of records if empty, key is taken from topmost level
in	<i>index</i>	- index in array of records
in	<i>key</i>	- key of key-value pair
in	<i>default_value</i>	- value to return, if no such key is found

Returns

parameter value

5.1.3.7 jsonKeyExist()

```
static bool DG::jsonKeyExist (
    const json & json_params,
    const std::string & section,
    int index,
    const std::string & key ) [inline], [static]
```

Check if given key exists in JSON array

Parameters

in	<i>json_params</i>	- JSON array
in	<i>section</i>	- top section name, which contains array of records if empty, key is checked on topmost level
in	<i>index</i>	- index in array of records
in	<i>key</i>	- key of key-value pair

Returns

true if such key exists in array

5.1.3.8 jsonSetOptionalValue()

```
template<typename T >
void DG::jsonSetOptionalValue (
    json & json_params,
    const std::string & section,
    int index,
    const std::string & key,
    const T & value )
```

Set value to JSON array; if it is not present, skip assignment

Parameters

in	<i>json_params</i>	- JSON array
in	<i>section</i>	- top section name, which contains array of records if empty, key is set on topmost level
in	<i>index</i>	- index in array of records
in	<i>key</i>	- key of key-value pair
in	<i>value</i>	- value to set

5.1.3.9 labelDictionary()

```
json DG::labelDictionary (
    const std::string & server,
    const std::string & model_name )
```

Get model label dictionary

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	<i>model_name</i>	specifies the AI model. To obtain valid model name, either modelFind() or modelzooListGet() functions should be used.

Returns

JSON object containing model label dictionary

5.1.3.10 messagePrepare()

```
static std::string DG::messagePrepare (
    const json & in ) [static]
```

prepare client-server response message string from input json

Parameters

in	server response json
----	----------------------

Returns

client response string

5.1.3.11 messagePrepareJson()

```
static json DG::messagePrepareJson (
    const json & in ) [static]
```

Prepare client-server response message json from input json

Parameters

<i>in</i>	server response json
-----------	----------------------

Returns

client response json

5.1.3.12 modelFind()

```
ModelInfo DG::modelFind (
    const std::string & server,
    const ModelQuery & query )
```

Find a model on a given AI server which matches given model query.

Parameters

<i>in</i>	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
<i>in</i>	<i>query</i>	is the ModelQuery structure of model parameters to look for.

Returns

model descriptor to be used in [AIModel](#) and [AIModelAsync](#) constructors. If the model is not found, the model descriptor contains an empty model name.

5.1.3.13 modelzooListGet()

```
void DG::modelzooListGet (
    const std::string & server,
    std::vector< ModelInfo > & modelzoo_list )
```

Get a list of supported models from AI server. In case of server connection errors throws std::exception.

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
out	<i>modelzoo_list</i>	is the output vector of ModelInfo structures; each structure contains essential model attributes such as model ID, model name, input width and input height. The ModelInfo structure extracted from this vector can be used as a parameter for construction of AIModel and AIModelAsync instances.

5.1.3.14 modelZooManage()

```
json DG::modelZooManage (
    const std::string & server,
    const json & req )
```

AI server model zoo management

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	<i>req</i>	- management request

Returns

results of management request completion (request-specific)

5.1.3.15 serverPing()

```
bool DG::serverPing (
    const std::string & server )
```

Ping server with an instantaneous command

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778. return true if no error occurred during the ping
----	---------------	---

5.1.3.16 shutdown()

```
void DG::shutdown (
    const std::string & server )
```

Send shutdown request to AI server [in] server is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.

5.1.3.17 systemInfo()

```
json DG::systemInfo (
    const std::string & server )
```

Return host system information dictionary.

Parameters

in	server	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
----	--------	--

5.1.3.18 traceManage()

```
json DG::traceManage (
    const std::string & server,
    const json & req )
```

AI server tracing facility management

Parameters

in	server	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	req	- management request

Returns

results of management request completion (request-specific)

5.1.3.19 versionGet()

```
std::string DG::versionGet ( )
```

Get the version of the library.

Returns

library version in the following format: MAJOR.MINOR.REVISION.GIT_REVISION

5.1.4 Variable Documentation

5.1.4.1 MODEL_PARAMS_CURRENT_VERSION

```
const int DG::MODEL_PARAMS_CURRENT_VERSION = 12
```

The most current version of JSON model configuration, supported by this version of software. Increment it each time you change any parameter definition or add/remove any parameter

5.1.4.2 MODEL_PARAMS_MIN_COMPATIBLE_VERSION

```
const int DG::MODEL_PARAMS_MIN_COMPATIBLE_VERSION = 1
```

The minimum compatible version of JSON model configuration, still supported by this version of software. Increase it when the software is modified such a way, that it stops supporting older JSON model configuration versions.

5.1.4.3 SECT_TOP

```
constexpr ModelParamsSection DG::SECT_TOP = { "", true } [constexpr]
```

JSON file top-level sections.

top-level section

Chapter 6

Class Documentation

6.1 DG::AIModel Class Reference

[AIModel](#) is DeGirum AI client API class for simple non-pipelined sequential inference.

```
#include <dg_model_api.h>
```

Public Member Functions

- [AIModel](#) (const std::string &server, const std::string &model_name, const [ModelParamsReadAccess](#) &model_params=[ModelParamsReadAccess](#){}, size_t connection_timeout_ms=[DEFAULT_CONNECTION_TIMEOUT_MS](#), size_t inference_timeout_ms=[DEFAULT_INFERENCE_TIMEOUT_MS](#))
- void [predict](#) (std::vector< std::vector< char > > &data, [json](#) &json_response)
- [~AIModel](#) ()

Destructor. Closes connection to AI server.

6.1.1 Detailed Description

[AIModel](#) is DeGirum AI client API class for simple non-pipelined sequential inference.

This class is used to perform AI model inference on AI server in a simple non-pipelined sequential manner. For more efficient (but somewhat more complex) asynchronous pipelined inference use [AIModelAsync](#) class.

On construction this class performs connection to AI server, selection of AI model, and setting model run-time parameters.

Once constructed, it can be used to perform sequential AI inference by invoking [predict\(\)](#) method. The [predict\(\)](#) methods accept the input frame data in various formats and return the inference results. These are blocking methods, i.e. they return execution only when the inference of an input frame is complete.

6.1.2 Constructor & Destructor Documentation

6.1.2.1 AIModel()

```
DG::AIModel::AIModel (
    const std::string & server,
    const std::string & model_name,
    const ModelParamsReadAccess & model_params = ModelParamsReadAccess({}),
    size_t connection_timeout_ms = DEFAULT_CONNECTION_TIMEOUT_MS,
    size_t inference_timeout_ms = DEFAULT_INFERENCE_TIMEOUT_MS ) [explicit]
```

Constructor. Performs connection to AI server, selection of AI model, and optionally setting model run-time parameters. In case of server connection errors throws std::exception.

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	<i>model_name</i>	specifies the AI model to be used for inference. To obtain valid model name, either modelFind() or modelzooListGet() functions should be used.
in	<i>model_params</i>	is runtime parameter collection, which defines the model runtime behavior. This is optional parameter: if not specified, then default runtime parameters (as defined in the model zoo) are used. ModelParamsWriter class instance can be used to conveniently define runtime parameters.
in	<i>connection_timeout_ms</i>	is the AI server connection timeout in milliseconds. This is optional parameter: by default it is set to 10 sec.
in	<i>inference_timeout_ms</i>	is the AI server inference timeout in milliseconds. This is optional parameter: by default it is set to 180 sec.

6.1.3 Member Function Documentation

6.1.3.1 predict()

```
void DG::AIModel::predict (
    std::vector< std::vector< char > > & data,
    json & json_response )
```

Run the AI inference on provided byte array. The byte array contains the frame data, which depends on selected frame format. It can be either JPEG or bitmap depending on the model parameters. In case of errors throws std::exception. This is blocking method, i.e. it returns execution only when the inference of an input frame is complete.

Parameters

in	<i>data</i>	is a vector of input data for each model input where each data element is a vector of bytes.
out	<i>json_response</i>	is the result of the inference. The response format depends on the model post-processor type.

The documentation for this class was generated from the following file:

- [dg_model_api.h](#)

6.2 DG::AIModelAsync Class Reference

[AIModelAsync](#) is DeGirum AI client API class for efficient pipelined asynchronous inference.

```
#include <dg_model_api.h>
```

Public Types

- using [callback_t](#) = std::function< void(const [json](#) &inference_result, const std::string &frame_info) >

Public Member Functions

- [AIModelAsync](#) (const std::string &server, const std::string &model_name, [callback_t](#) callback, const [ModelParamsReadAccess](#) &model_params=[ModelParamsReadAccess](#){}, size_t frame_queue_depth=[DEFAULT_FRAME_QUEUE_DEPTH](#), size_t connection_timeout_ms=[DEFAULT_CONNECTION_TIMEOUT_MS](#), size_t inference_timeout_ms=[DEFAULT_INFERENCE_TIMEOUT_MS](#))
- std::string [lastError](#) () const
- int [outstandingResultsCountGet](#) () const
Get the number of outstanding inference results posted so far.
- void [predict](#) (std::vector< std::vector< char > > &data, const std::string &frame_info="")
- void [setCallback](#) ([callback_t](#) callback)
- void [waitCompletion](#) ()
- [~AIModelAsync](#) ()

6.2.1 Detailed Description

[AIModelAsync](#) is DeGirum AI client API class for efficient pipelined asynchronous inference.

This class is used to perform AI model inference on AI server in efficient pipelined asynchronous manner using the mechanism of callbacks. For simple (but less efficient) synchronous non-pipelined inference use [AIModel](#) class.

On construction this class performs connection to AI server, selection of AI model, setting model run-time parameters, and installation of the client callback function. The client callback function is used to pass the inference results from the AI server to the client code.

Once constructed, it can be used to perform asynchronous AI inference by invoking [predict\(\)](#) method. The [predict\(\)](#) methods accept the input frame data and initiate the inference on the AI server. Each of those methods is a non-blocking method, i.e. it returns execution immediately after posting the frame data to the AI server. This allows calling those methods in a loop without waiting for the inference results, achieving the maximum AI server utilization. Once the inference of a frame is complete, and the inference result is received from the AI server, the client callback is invoked to dispatch the inference result. Such result handling via callback mechanism is performed in a thread, separate from the main execution thread. It means that the client callback function is called in asynchronous manner, thus the name of the class.

6.2.2 Member Typedef Documentation

6.2.2.1 callback_t

```
using DG::AIModelAsync::callback_t = std::function< void( const json &inference_result, const
std::string &frame_info ) >
```

User callback type. The callback is called asynchronously from the main execution thread as soon as prediction result is ready. Consecutive prediction result in a form of JSON array is passed as the `inference_result` argument. Corresponding frame info string (provided to `predict()` call) is passed as the `frame_info` argument.

6.2.3 Constructor & Destructor Documentation

6.2.3.1 AIModelAsync()

```
DG::AIModelAsync::AIModelAsync (
    const std::string & server,
    const std::string & model_name,
    callback_t callback,
    const ModelParamsReadAccess & model_params = ModelParamsReadAccess({}),
    size_t frame_queue_depth = DEFAULT_FRAME_QUEUE_DEPTH,
    size_t connection_timeout_ms = DEFAULT_CONNECTION_TIMEOUT_MS,
    size_t inference_timeout_ms = DEFAULT_INFERENCE_TIMEOUT_MS ) [explicit]
```

Constructor. Performs connection to AI server, selection of AI model, installing client callback, and optionally setting model run-time parameters. In case of server connection errors throws `std::exception`.

Parameters

in	<i>server</i>	is a string specifying server domain name/IP address and port. Format: "domain_name:port" or "xxx.xxx.xxx.xxx:port". If port is omitted, the default port is 8778.
in	<i>model_name</i>	specifies the AI model to be used for inference. To obtain valid model name, either <code>modelFind()</code> or <code>modelzooListGet()</code> functions should be used.
in	<i>callback</i>	is user callback functional, which will be called asynchronously from the main execution thread as soon as prediction result is ready.
in	<i>model_params</i>	is runtime parameter collection, which defines the model runtime behavior. This is optional parameter: if not specified, then default runtime parameters (as defined in the model zoo) are used. <code>ModelParamsWriter</code> class instance can be used to conveniently define runtime parameters.
in	<i>frame_queue_depth</i>	is the depth of the internal frame queue. If <code>predict()</code> methods are invoked too often and the number of non-processed (aka "outstanding") frames exceeds this parameter, the consecutive call to any <code>predict()</code> method will be blocked until the number of outstanding frames in the queue becomes smaller than the queue depth thus allowing to post one more frame. This is optional parameter: by default it is set to 8 frames.
in	<i>connection_timeout_ms</i>	is the AI server connection timeout in milliseconds. This is optional parameter: by default it is set to 10 sec.
in	<i>inference_timeout_ms</i>	is the AI server inference timeout in milliseconds. This is optional parameter: by default it is set to 180 sec.

6.2.3.2 ~AIModelAsync()

```
DG::AIModelAsync::~~AIModelAsync ( )
```

Destructor. Waits until all outstanding results are received and then closes the connection to AI server. Note: in case of server runtime error, all frames posted after that error was detected, will not be processed.

6.2.4 Member Function Documentation

6.2.4.1 lastError()

```
std::string DG::AIModelAsync::lastError ( ) const
```

If ever during consecutive calls to [predict\(\)](#) methods AI server reported a run-time error, then this method will return the error message string, otherwise it returns an empty string. Note: in case of server runtime error, all frames posted after that error was detected, will not be processed.

6.2.4.2 predict()

```
void DG::AIModelAsync::predict (
    std::vector< std::vector< char > > & data,
    const std::string & frame_info = "" )
```

Start the inference on given byte data vector. The byte vector contains the frame data, which depends on selected frame format. It can be either JPEG or bitmap depending on the model parameters. In case of errors throws `std::exception`. This is non-blocking call meaning that it returns execution immediately after posting the frame data to the AI server.

Parameters

in	<i>data</i>	is a vector of input data for each model input where each data element is a vector of bytes.
in	<i>frame_info</i>	is optional frame information string to be passed to the client callback along with the frame result. You can pass arbitrary information as frame info. This simplifies matching results to frames in the client callback.

6.2.4.3 setCallback()

```
void DG::AIModelAsync::setCallback (
    callback_t callback )
```

Set user callback

Parameters

<code>in</code>	<code>callback</code>	is user callback functional, which will be called asynchronously from the main execution thread as soon as prediction result is ready.
-----------------	-----------------------	--

6.2.4.4 waitCompletion()

```
void DG::AIModelAsync::waitCompletion ( )
```

Wait for completion of all outstanding inferences. This is blocking call: it returns when all outstanding frames are processed by AI server and all results are dispatched via client callback. You can continue calling [predict\(\)](#) after call to this method - frame processing will restart automatically.

The documentation for this class was generated from the following file:

- [dg_model_api.h](#)

6.3 DG::BasicTensor Class Reference

```
#include <dg_tensor_structs.h>
```

Classes

- struct [quant_params_t](#)
Quantization parameters (how to convert from integer/quantized data back to floating point)

Public Types

- using [shape_t](#) = std::vector< size_t >
Tensor shape vector type.

Public Member Functions

- template<class CONTAINER >
auto [adopt](#) (CONTAINER &&source, int32_t id=0, const std::string &name="", const [shape_t](#) &shape={}, const [quant_params_t](#) &quant_params={}) -> decltype(source.data(), source.size(), void())
- template<class CONTAINER >
auto [alloc](#) (const CONTAINER &source, bool do_copy, int32_t id=0, const std::string &name="", const [shape_t](#) &shape={}, const [quant_params_t](#) &quant_params={}) -> decltype(source.data(), source.size(), void())
- template<typename T >
void [alloc](#) (int32_t id, const std::string &name, const [shape_t](#) &shape, const [quant_params_t](#) &quant_params={}, T *ext_lin_buffer=nullptr)
- void [alloc](#) (int32_t id, const std::string &name, const [shape_t](#) &shape, DGType data_type, const [quant_params_t](#) &quant_params={}, void *ext_lin_buffer=nullptr)
- [BasicTensor](#) () noexcept

- Default constructor.*
- `BasicTensor (BasicTensor &&move) noexcept`
- Move constructor.*
- `BasicTensor (const BasicTensor &)=delete`
- Deleted copy constructor.*
- `template<class CONTAINER >`
`BasicTensor (const CONTAINER &source, bool do_copy=true, int32_t id=0, const std::string &name="", const shape_t &shape={}, const quant_params_t &quant_params={})`
- `template<typename T >`
`BasicTensor (int32_t id, const std::string &name, const shape_t &shape, const quant_params_t &quant_params={}, T *ext_lin_buffer=nullptr)`
- `BasicTensor (int32_t id, const std::string &name, const shape_t &shape, DGType data_type, const quant_params_t &quant_params={}, void *ext_lin_buffer=nullptr)`
- `BasicTensor clone () const`
 - Create and return exact copy of itself.*
- `template<typename T >`
`BasicTensor convert () const`
- `BasicTensor convert (DGType to_type) const`
- `BasicTensor copy () const`
- `template<typename T >`
`T * data ()`
- `template<typename T >`
`const T * data () const`
- `DGType dataTypeGet () const`
 - Get DG data type of tensor element.*
- `void dealloc ()`
 - Deallocate tensor data and clear tensor.*
- `template<typename T_IN, typename T_OUT >`
`void dequantize ()`
- `size_t elementSizeGet () const`
 - Get tensor element size in bytes.*
- `bool empty () const`
 - Check if tensor is empty.*
- `int32_t id () const`
 - Get id.*
- `FloatCompareResult< double > isEqualData (const BasicTensor &rhs, double maxRelDiff) const`
- `template<typename T >`
`FloatCompareResult< double > isEqualData (const std::vector< T > &rhs, double maxRelDiff) const`
- `bool isEqualDataShape (const BasicTensor &rhs) const`
- `bool isEqualStruct (const BasicTensor &rhs) const`
- `bool isExternal () const`
 - Is external buffer?*
- `bool isNull () const`
 - Check if tensor is null.*
- `size_t linearSizeGet () const`
 - Get tensor linear size in elements of current type.*
- `size_t linearSizeGet_bytes () const`
 - Get tensor linear size in bytes.*
- `const std::string & name () const`
 - Get name.*
- `void NCHWToNHWC ()`
 - Wrapper function for NCHWToNHWC<T>()*

- template<typename T >
void [NCHWToNHWC](#) ()
Roll NCHW -> NHWC.
- void [NHWCToNCHW](#) ()
Wrapper function for [NHWCToNCHW<T>\(\)](#)
- template<typename T >
void [NHWCToNCHW](#) ()
Roll NHWC -> NCHW.
- std::string [numpyTypeGet](#) () const
Return numpy-compatible type string.
- [BasicTensor](#) & operator= (BasicTensor &&move)
Move assignment operator.
- [BasicTensor](#) & operator= (const [BasicTensor](#) &)=delete
Deleted copy assignment operator.
- template<typename T_IN , typename T_OUT >
void [quantize](#) ()
- const [quant_params_t](#) & [quantParams](#) () const
Get quantization parameters.
- void [quantParamsSet](#) (const [quant_params_t](#) &qp)
- void [reinterpretShapeForYolo](#) ()
- void [reinterpretShapeScaled](#) (size_t last_dim, float32_t scale)
- void [reshapeCombinedDims](#) (const std::vector< size_t > dims_to_combine)
- void [reshapeTo](#) (const [shape_t](#) &new_shape, DGType new_type=DG_UNDEFINED)
- void [reshapeTo](#) (size_t dim)
- void [reshapeToNHWC](#) ()
Reshape tensor to 4-D NHWC shape.
- const [shape_t](#) & [shape](#) () const
Get shape vector.
- std::string [shapeStringGet](#) () const
- const std::type_info & [typeInfo](#) () const
Return type info.
- template<typename T = void>
T * [untypedData](#) ()
Return raw data buffer.
- template<typename T = void>
T * [untypedData](#) () const
Return raw data buffer.
- [~BasicTensor](#) ()
Destructor.

Static Public Member Functions

- static size_t [linearSizeCalc](#) (const [shape_t](#) &shape)
- static std::string [numpyTypeGet](#) (const std::string &dg_type)
Return numpy-compatible type string from given [DG](#) type string.
- static std::string [shapeStringGet](#) (const [shape_t](#) &shape)

6.3.1 Detailed Description

Basic tensor container class. Based on "linear buffer with dimension array" approach. Dynamically-typed. Supports both internally or externally allocated buffers.

6.3.2 Constructor & Destructor Documentation

6.3.2.1 BasicTensor() [1/3]

```
template<typename T >
DG::BasicTensor::BasicTensor (
    int32_t id,
    const std::string & name,
    const shape_t & shape,
    const quant_params_t & quant_params = {},
    T * ext_lin_buffer = nullptr ) [inline]
```

Constructor: allocate or assign tensor memory according to dimension vector (static typing version)

Parameters

in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>quant_params</i>	- tensor quantization parameters
in	<i>ext_lin_buffer</i>	- optional external linear buffer pointer; if nullptr, buffer will be allocated internally

6.3.2.2 BasicTensor() [2/3]

```
DG::BasicTensor::BasicTensor (
    int32_t id,
    const std::string & name,
    const shape_t & shape,
    DGType data_type,
    const quant_params_t & quant_params = {},
    void * ext_lin_buffer = nullptr ) [inline]
```

Constructor: allocate or assign tensor memory according to dimension vector (dynamic typing version)

Parameters

in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>data_type</i>	- tensor element data type id
in	<i>quant_params</i>	- tensor quantization parameters
in	<i>ext_lin_buffer</i>	- optional external linear buffer pointer; if nullptr, buffer will be allocated internally

6.3.2.3 BasicTensor() [3/3]

```
template<class CONTAINER >
DG::BasicTensor::BasicTensor (
    const CONTAINER & source,
    bool do_copy = true,
    int32_t id = 0,
    const std::string & name = "",
    const shape_t & shape = {},
    const quant_params_t & quant_params = {} ) [inline], [explicit]
```

Constructor: allocate tensor memory and copy data from given linear container. CONTAINER type should have value_type, data(), and size() methods. If shape is empty, 1-D tensor the size of the linear container will be allocated. The size of the data copied is the minimum of the linear container size and the shape dimensions product.

Parameters

in	<i>source</i>	- linear container (such as string or vector) with tensor data
in	<i>do_copy</i>	- when true, copy data from container into tensor-owned buffer, otherwise assign pointer to container data
in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>quant_params</i>	- tensor quantization parameters

6.3.3 Member Function Documentation

6.3.3.1 adopt()

```
template<class CONTAINER >
auto DG::BasicTensor::adopt (
    CONTAINER && source,
    int32_t id = 0,
    const std::string & name = "",
    const shape_t & shape = {},
    const quant_params_t & quant_params = {} ) -> decltype( source.data(), source.↔
size(), void() )

[inline]
```

Adopt data from given linear container by moving it to internal storage. CONTAINER type should have value_type, data(), and size() methods.

Parameters

in	<i>source</i>	- linear container (such as string or vector) with tensor data
in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>quant_params</i>	- tensor quantization parameters

6.3.3.2 alloc() [1/3]

```
template<class CONTAINER >
auto DG::BasicTensor::alloc (
    const CONTAINER & source,
    bool do_copy,
    int32_t id = 0,
    const std::string & name = "",
    const shape_t & shape = {},
    const quant_params_t & quant_params = {} ) -> decltype( source.data(), source.size(), void() )
```

[inline]

Allocate tensor memory and copy data from given linear container. CONTAINER type should have value_type, data(), and size() methods. If shape is empty, 1-D tensor the size of the linear container will be allocated. The size of the data copied is the minimum of the linear container size and the shape dimensions product.

Parameters

in	<i>source</i>	- linear container (such as string or vector) with tensor data
in	<i>do_copy</i>	- when true, copy data from container into tensor-owned buffer, otherwise assign pointer to container data
in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>quant_params</i>	- tensor quantization parameters

6.3.3.3 alloc() [2/3]

```
template<typename T >
void DG::BasicTensor::alloc (
    int32_t id,
    const std::string & name,
    const shape_t & shape,
    const quant_params_t & quant_params = {},
    T * ext_lin_buffer = nullptr ) [inline]
```

Allocate or assign tensor memory according to dimension vector (static typing version)

Parameters

in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>quant_params</i>	- tensor quantization parameters
in	<i>ext_lin_buffer</i>	- optional external linear buffer pointer; if nullptr, buffer will be allocated internally

6.3.3.4 alloc() [3/3]

```
void DG::BasicTensor::alloc (
    int32_t id,
    const std::string & name,
    const shape_t & shape,
    DGType data_type,
    const quant_params_t & quant_params = {},
    void * ext_lin_buffer = nullptr ) [inline]
```

Allocate or assign tensor memory according to dimension vector (dynamic typing version)

Parameters

in	<i>id</i>	- tensor ID
in	<i>name</i>	- tensor name
in	<i>shape</i>	- tensor shape vector, which defines dimensions of new tensor
in	<i>data_type</i>	- tensor element data type id
in	<i>quant_params</i>	- tensor quantization parameters
in	<i>ext_lin_buffer</i>	- optional external linear buffer pointer; if nullptr, buffer will be allocated internally

6.3.3.5 convert() [1/2]

```
template<typename T >
BasicTensor DG::BasicTensor::convert ( ) const [inline]
```

Convert tensor to new element type with casting and copying data

Template Parameters

<i>output</i>	data type T
---------------	-------------

Returns

new owning tensor of type T

6.3.3.6 convert() [2/2]

```
BasicTensor DG::BasicTensor::convert (
    DGType to_type ) const [inline]
```

Convert tensor to new element type with casting and copying data

Parameters

in	<i>to_type</i>	- new element type
----	----------------	--------------------

Returns

new owning tensor

6.3.3.7 copy()

```
BasicTensor DG::BasicTensor::copy ( ) const [inline]
```

Create and return deep copy of itself. External buffer data will be duplicated.

6.3.3.8 data() [1/2]

```
template<typename T >
T* DG::BasicTensor::data ( ) [inline]
```

Return non-const typed pointer to underlying linear buffer. Returns nullptr if template type does not match actual runtime tensor element type

6.3.3.9 data() [2/2]

```
template<typename T >
const T* DG::BasicTensor::data ( ) const [inline]
```

Return const typed pointer to underlying linear buffer. Returns nullptr if template type does not match actual runtime tensor element type

6.3.3.10 dequantize()

```
template<typename T_IN , typename T_OUT >
void DG::BasicTensor::dequantize ( ) [inline]
```

Dequantize tensor according to current quantization settings from type T_IN to type T_OUT New linear buffer will be internally allocated to receive dequantized data.

6.3.3.11 isEqualData() [1/2]

```
FloatCompareResult< double > DG::BasicTensor::isEqualData (
    const BasicTensor & rhs,
    double maxRelDiff ) const [inline]
```

Check, if this object and given object have equal data contents

Parameters

in	<i>rhs</i>	- object to compare with
in	<i>maxRelDiff</i>	- maximum relative difference for floating point comparison (ignored for integer data)

6.3.3.12 isEqualData() [2/2]

```
template<typename T >
FloatCompareResult< double > DG::BasicTensor::isEqualData (
    const std::vector< T > & rhs,
    double maxRelDiff ) const [inline]
```

Check, if this object and given vector have equal data contents

Parameters

in	<i>rhs</i>	- vector to compare with
in	<i>maxRelDiff</i>	- maximum relative difference for floating point comparison (ignored for integer data)

6.3.3.13 isEqualDataShape()

```
bool DG::BasicTensor::isEqualDataShape (
    const BasicTensor & rhs ) const [inline]
```

Check, if this object and given object have equal shape, element size and number of elements

Parameters

in	<i>rhs</i>	- object to compare with
----	------------	--------------------------

6.3.3.14 isEqualStruct()

```
bool DG::BasicTensor::isEqualStruct (
    const BasicTensor & rhs ) const [inline]
```

Check, if this object and given object have equal structure

Parameters

in	<i>rhs</i>	- object to compare with
----	------------	--------------------------

6.3.3.15 linearSizeCalc()

```
static size_t DG::BasicTensor::linearSizeCalc (
    const shape_t & shape ) [inline], [static]
```

Calculate tensor linear size by shape

Parameters

in	<i>shape</i>	- tensor shape vector
----	--------------	-----------------------

Returns

linear size in elements of current type

6.3.3.16 quantize()

```
template<typename T_IN , typename T_OUT >
void DG::BasicTensor::quantize ( ) [inline]
```

Quantize tensor according to current quantization settings from type T_IN to type T_OUT New linear buffer will be internally allocated to receive quantized data.

6.3.3.17 quantParamsSet()

```
void DG::BasicTensor::quantParamsSet (
    const quant_params_t & qp ) [inline]
```

Set quantization parameters

Parameters

in	<i>qp</i>	- new quantization parameters
----	-----------	-------------------------------

6.3.3.18 reinterpretShapeForYolo()

```
void DG::BasicTensor::reinterpretShapeForYolo ( ) [inline]
```

Reinterpret shape tensor to 3d to be used for YOLO detection/pose/segmentation models The last dimension is interpreted as channels, and all other dimensions are flattened 1xHxWxC -> 1x(H*W)xC HxWx1xC -> 1x(H*W)xC

6.3.3.19 reinterpretShapeScaled()

```
void DG::BasicTensor::reinterpretShapeScaled (
    size_t last_dim,
    float32_t scale ) [inline]
```

If the 3rd dimension is equal to last_dim, reshape tensor by scaling up 2nd dimension by scale and scaling down 3rd dimension by scale. E.g. If last_dim = 255 and scale = 3, then shape (1 x 6400 x 255) would be reshaped as 1 x 6400 x 255 -> 1 x 19200 x 85

6.3.3.20 reshapeCombineDims()

```
void DG::BasicTensor::reshapeCombineDims (
    const std::vector< size_t > dims_to_combine ) [inline]
```

Reshape tensor of multiple dimensions to a smaller number of dimensions by combining specified dimensions. E.g. reshapeCombineDims({0, 1, 2}) will combine first 3 dimensions into one.

6.3.3.21 reshapeTo() [1/2]

```
void DG::BasicTensor::reshapeTo (
    const shape_t & new_shape,
    DGType new_type = DG_UNDEFINED ) [inline]
```

Reshape tensor to given shape and element type

Parameters

in	<i>new_shape</i>	- new shape vector
in	<i>new_type</i>	- new element type, if DG_UNDEFINED, use current type Linear size after reshape must be the same as before reshape.

6.3.3.22 reshapeTo() [2/2]

```
void DG::BasicTensor::reshapeTo (
    size_t dim ) [inline]
```

Reshape tensor to given shape When new shape is bigger than current, shape is appended with unity values When new shape is smaller than current, shape is reduced, and last shape dimension is multiplied by product of reduced dimensions

6.3.3.23 shapeStringGet() [1/2]

```
std::string DG::BasicTensor::shapeStringGet ( ) const [inline]
```

Construct shape string in "dim1xdim2x..." format

Returns

shape string

6.3.3.24 shapeStringGet() [2/2]

```
static std::string DG::BasicTensor::shapeStringGet (
    const shape_t & shape ) [inline], [static]
```

Construct shape string in "dim1xdim2x..." format

Parameters

<code>in</code>	<code>shape</code>	- tensor shape vector
-----------------	--------------------	-----------------------

Returns

shape string

The documentation for this class was generated from the following file:

- [dg_tensor_structs.h](#)

6.4 DG::JsonHelper Class Reference

JSON helper class: contains assorted static methods to operate with JSON configurations.

```
#include <dg_json_helpers.h>
```

Public Types

- using [serial_container_t](#) = json::binary_t::container_type
Container type for serialized JSON array (it is std::vector< unit8_t >)

Static Public Member Functions

- static std::string [errorCheck](#) (const [json](#) &response, const std::string &source, bool do_throw=true)
- static [json](#) [jsonDeserialize](#) (const [serial_container_t](#) &v)
- static [json](#) [jsonDeserializeStr](#) (const std::string &v)
- static [serial_container_t](#) [jsonSerialize](#) (const [json](#) &j)
- static std::string [jsonSerializeStr](#) (const [json](#) &j)
- static [json](#) [parse](#) (const std::string &json_cfg, const char *file, const char *line, const char *func)
- static [json](#) [parse_ignore_errors](#) (const std::string json_cfg)
- static [BasicTensor](#) [tensorDeserialize](#) (const [json](#) &j)
- static [json](#) [tensorSerialize](#) (const [BasicTensor](#) &t)

6.4.1 Detailed Description

JSON helper class: contains assorted static methods to operate with JSON configurations.

6.4.2 Member Function Documentation

6.4.2.1 errorCheck()

```
static std::string DG::JsonHelper::errorCheck (
    const json & response,
    const std::string & source,
    bool do_throw = true ) [inline], [static]
```

Check server JSON response for errors and throw exception, if any

Parameters

in	<i>response</i>	- JSON response from server/core
in	<i>source</i>	- description of the server command initiator
in	<i>do_throw</i>	- if true, throws exception in case of error

Returns

original error string, if error is detected, empty string otherwise

6.4.2.2 jsonDeserialize()

```
static json DG::JsonHelper::jsonDeserialize (
    const serial_container_t & v ) [inline], [static]
```

Deserialize given byte vector with msgpack data into JSON array

Parameters

in	<i>v</i>	- byte vector with msgpack data to deserialize
----	----------	--

Returns

JSON array

6.4.2.3 jsonDeserializeStr()

```
static json DG::JsonHelper::jsonDeserializeStr (
    const std::string & v ) [inline], [static]
```

Deserialize given string with msgpack data into JSON array

Parameters

in	<i>v</i>	- byte vector with msgpack data to deserialize
----	----------	--

Returns

JSON array

6.4.2.4 jsonSerialize()

```
static serial_container_t DG::JsonHelper::jsonSerialize (
    const json & j ) [inline], [static]
```

Serialize given JSON array into byte vector using conversion to msgpack

Parameters

in	<i>j</i>	- JSON array to serialize
----	----------	---------------------------

Returns

byte vector with JSON array contents converted to msgpack representation

6.4.2.5 jsonSerializeStr()

```
static std::string DG::JsonHelper::jsonSerializeStr (
    const json & j ) [inline], [static]
```

Serialize given JSON array into string using conversion to msgpack

Parameters

in	<i>j</i>	- JSON array to serialize
----	----------	---------------------------

Returns

string with JSON array contents converted to msgpack representation

6.4.2.6 parse()

```
static json DG::JsonHelper::parse (
    const std::string & json_cfg,
    const char * file,
    const char * line,
    const char * func ) [inline], [static]
```

Parse given JSON string with proper exception handling

Parameters

in	<i>json_cfg</i>	- JSON string to parse
in	<i>file</i>	- file where parsing happens
in	<i>line</i>	- line in file where parsing happens
in	<i>func</i>	- function where parsing happens

Returns

JSON array

6.4.2.7 parse_ignore_errors()

```
static json DG::JsonHelper::parse_ignore_errors (
    const std::string json_cfg ) [inline], [static]
```

Parse given JSON string with swallowing exceptions

Parameters

in	<i>json_cfg</i>	- JSON string to parse
----	-----------------	------------------------

Returns

JSON array or empty array, if parsing failed

6.4.2.8 tensorDeserialize()

```
static BasicTensor DG::JsonHelper::tensorDeserialize (
    const json & j ) [inline], [static]
```

Convert JSON array to basic tensor

Parameters

in	<i>j</i>	- JSON array
----	----------	--------------

Returns

basic tensor

6.4.2.9 tensorSerialize()

```
static json DG::JsonHelper::tensorSerialize (
    const BasicTensor & t ) [inline], [static]
```

Convert basic tensor to JSON array

Parameters

in	<i>t</i>	- basic tensor
----	----------	----------------

Returns

JSON array

The documentation for this class was generated from the following file:


- [dg_json_helpers.h](#)

6.5 DG::ModelInfo Struct Reference

[ModelInfo](#) is the model identification structure. It keeps AI model key attributes.

```
#include <dg_client_structs.h>
```

Collaboration diagram for DG::ModelInfo:



Public Attributes

- [DG::ModelParamsWriter](#) `extended_params`
extended model parameters
- `std::string` [name](#)
model string name

6.5.1 Detailed Description

[ModelInfo](#) is the model identification structure. It keeps AI model key attributes.

The documentation for this struct was generated from the following file:


- [dg_client_structs.h](#)

6.6 DG::ModelParams< Base, > Class Template Reference


[ModelParams](#) is model parameters collection with user-defined access rights.

```
#include <dg_model_parameters.h>
```

Inheritance diagram for DG::ModelParams< Base, >:



Collaboration diagram for DG::ModelParams< Base, >:



Public Member Functions

- [ModelParams](#) (const char *json_text="{}")
- [ModelParams](#) (const [json](#) &json_cfg)
- [ModelParams](#) (const std::string &json_text)
- [ModelParams](#) ([json](#) &&json_cfg)
- [ModelParams](#) ([ModelParams](#) &&rhs)
Move constructor.
- [ModelParams](#) ([ModelParams](#) const &rhs)
Copy constructor.
- [ModelParams](#) & [operator=](#) ([ModelParams](#) &&rhs)
Move assignment.
- [ModelParams](#) & [operator=](#) ([ModelParams](#) const &rhs)
Copy assignment.

Additional Inherited Members

6.6.1 Detailed Description

```
template<class Base = ModelParamsReadAccess, std::enable_if_t< std::is_base_of_v< ModelParamsReadAccess, Base >,
bool > = false>
class DG::ModelParams< Base, >
```

[ModelParams](#) is model parameters collection with user-defined access rights.

This class provides programmatic type-safe access to model parameters defined in JSON model configuration while owning that JSON array. Access (read or write) is defined by the Base template parameter, which is used as the base class.

Template Parameters

<i>Base</i>	- ModelParamsReadAccess-derived base class: it can be ModelParamsReadAccess or ModelParamsWriteAccess
-------------	---

6.6.2 Constructor & Destructor Documentation

6.6.2.1 ModelParams() [1/4]

```
template<class Base = ModelParamsReadAccess, std::enable_if_t< std::is_base_of_v< ModelParamsReadAccess, Base >, bool > = false>
DG::ModelParams< Base, >::ModelParams (
    const char * json_text = "{}" ) [inline]
```

Constructor. Creates model parameter collection by parsing JSON text in character array

Parameters

in	<i>json_text</i>	- JSON text with model parameters. It will be parsed and stored in internal JSON array.
----	------------------	---

6.6.2.2 ModelParams() [2/4]

```
template<class Base = ModelParamsReadAccess, std::enable_if_t< std::is_base_of_v< Model↵
ParamsReadAccess, Base >, bool > = false>
DG::ModelParams< Base, >::ModelParams (
    const std::string & json_text ) [inline], [explicit]
```

Constructor. Creates model parameter collection by parsing JSON text in std::string

Parameters

in	<i>json_text</i>	- JSON text with model parameters. It will be parsed and stored in internal JSON array.
----	------------------	---

6.6.2.3 ModelParams() [3/4]

```
template<class Base = ModelParamsReadAccess, std::enable_if_t< std::is_base_of_v< Model↵
ParamsReadAccess, Base >, bool > = false>
DG::ModelParams< Base, >::ModelParams (
    const json & json_cfg ) [inline], [explicit]
```

Constructor. Creates model parameter collection from given JSON array

Parameters

in	<i>json_cfg</i>	- JSON array with model parameters. It will be copied into internal JSON array.
----	-----------------	---

6.6.2.4 ModelParams() [4/4]

```
template<class Base = ModelParamsReadAccess, std::enable_if_t< std::is_base_of_v< Model↵
ParamsReadAccess, Base >, bool > = false>
DG::ModelParams< Base, >::ModelParams (
    json && json_cfg ) [inline], [explicit]
```

Constructor with move semantic. Creates model parameter collection from given JSON array

Parameters

in	<i>json_cfg</i>	- JSON array with model parameters. It will be moved into internal JSON array.
----	-----------------	--

The documentation for this class was generated from the following file:


- [dg_model_parameters.h](#)

6.7 DG::ModelParamsReadAccess Class Reference

[ModelParamsReadAccess](#) is read-only accessor to model parameters.

```
#include <dg_model_parameters.h>
```

Inheritance diagram for DG::ModelParamsReadAccess:



Public Member Functions

- [json fullJsonGet](#) () const
- const [json & jsonGet](#) () const
- size_t [modellInputSizeGet](#) () const
Helper method: get model input size.
- [ModelParamsReadAccess](#) (const [json](#) &config)
- std::vector< size_t > [modelShapeGet](#) (size_t inp_idx, size_t expected_size)
- template<typename T >
T [None](#) (size_t idx=0) const
Stub for empty getter fallback.
- bool [None_exist](#) (size_t idx=0) const
Stub for empty existence checker fallback.
- template<typename T >
T [None_get](#) (size_t idx=0) const
Stub for empty getter fallback.
- operator const [json](#) & () const
- operator std::string () const
- const [json](#) & [sectionGet](#) (const std::string §ion, size_t idx=0) const
- size_t [sectionHashGet](#) (const std::string §ion, size_t idx=0) const
- size_t [sectionSizeGet](#) (const std::string §ion) const

Protected Member Functions

- bool [paramExist](#) (const char *section, const char *key, size_t idx, bool(ModelParamsReadAccess::*fallback)(size_t) const) const
- template<typename T >
T [paramGet](#) (const char *section, const char *key, bool is_mandatory, const T &default_value, size_t idx, T(ModelParamsReadAccess::*fallback)(size_t) const) const

Protected Attributes

- const [json](#) & [m_cfg_ro](#)
non-owning reference to JSON array with model configuration

6.7.1 Detailed Description

[ModelParamsReadAccess](#) is read-only accessor to model parameters.

This class provides programmatic type-safe read access to model parameters defined in JSON model configuration. It keeps non-owning const reference to underlying JSON array. For each model parameter it provides getter method, which name matches the parameter name as it appears in JSON array.

6.7.2 Constructor & Destructor Documentation

6.7.2.1 ModelParamsReadAccess()

```
DG::ModelParamsReadAccess::ModelParamsReadAccess (
    const json & config ) [inline]
```

Constructor. Attaches model parameter read-only accessor to JSON array

Parameters

in	<i>config</i>	- JSON array with model configuration
--------------------	---------------	---------------------------------------

6.7.3 Member Function Documentation

6.7.3.1 fullJsonGet()

```
json DG::ModelParamsReadAccess::fullJsonGet ( ) const [inline]
```

Function to return the full JSON of model parameters This returns the model parameter JSON merged with a JSON of default values for all the parameters that aren't listed

6.7.3.2 jsonGet()

```
const json& DG::ModelParamsReadAccess::jsonGet ( ) const [inline]
```

Access to underlying JSON array

Returns

constant non-owning reference to JSON array with model configuration

6.7.3.3 modelShapeGet()

```
std::vector< size_t > DG::ModelParamsReadAccess::modelShapeGet (
    size_t inp_idx,
    size_t expected_size ) [inline]
```

Helper method: get model input shape

Parameters

in	<i>inp_idx</i>	- input index
in	<i>expected_size</i>	- expected size of the shape vector, 0 if not checked

Returns

model input shape vector

6.7.3.4 operator const json &()

```
DG::ModelParamsReadAccess::operator const json & ( ) const [inline]
```

Access to underlying JSON array

Returns

constant non-owning reference to JSON array with model configuration

6.7.3.5 operator std::string()

```
DG::ModelParamsReadAccess::operator std::string ( ) const [inline]
```

Access to underlying JSON array as string

Returns

JSON text with model configuration

6.7.3.6 paramExist()

```
bool DG::ModelParamsReadAccess::paramExist (
    const char * section,
    const char * key,
    size_t idx,
    bool (ModelParamsReadAccess::*)(size_t) const fallback ) const [inline], [protected]
```

Check parameter for existence

Template Parameters

<i>T</i>	- parameter type
----------	------------------

Parameters

in	<i>section</i>	- top section name; if empty, key is taken from topmost level
in	<i>key</i>	- key of key-value pair
in	<i>idx</i>	- array index inside section object
in	<i>fallback</i>	- existence checker for fallback parameter

Returns

parameter value

6.7.3.7 paramGet()

```
template<typename T >
T DG::ModelParamsReadAccess::paramGet (
    const char * section,
    const char * key,
    bool is_mandatory,
    const T & default_value,
    size_t idx,
    T (ModelParamsReadAccess::*)(size_t) const fallback ) const [inline], [protected]
```

Get parameter from JSON array

Template Parameters

<i>T</i>	- parameter type
----------	------------------

Parameters

in	<i>section</i>	- top section name; if empty, key is taken from topmost level
in	<i>key</i>	- key of key-value pair
in	<i>is_mandatory</i>	- if parameter mandatory
in	<i>default_value</i>	- value to return, if no such key is found

Parameters

in	<i>idx</i>	- array index inside section object
in	<i>fallback</i>	- fallback parameter getter

Returns

parameter value

6.7.3.8 sectionGet()

```
const json& DG::ModelParamsReadAccess::sectionGet (
    const std::string & section,
    size_t idx = 0 ) const [inline]
```

Access to underlying JSON sub-array for particular section

Parameters

in	<i>section</i>	- section name to access (one of SECT_xxx constants)
in	<i>idx</i>	- array index inside section object

Returns

constant non-owning reference to JSON array with particular model section configuration

6.7.3.9 sectionHashGet()

```
size_t DG::ModelParamsReadAccess::sectionHashGet (
    const std::string & section,
    size_t idx = 0 ) const [inline]
```

Compute the hash of the parameter values in a given section

Parameters

in	<i>section</i>	- section name to compute hash for
in	<i>idx</i>	- array index inside section object

Returns

std::hash for a given section

6.7.3.10 sectionSizeGet()

```
size_t DG::ModelParamsReadAccess::sectionSizeGet (
    const std::string & section ) const [inline]
```

Get size of underlying JSON sub-array for particular section

Parameters

in	<i>section</i>	- section name to access (one of SECT_xxx constants)
----	----------------	--

Returns

section sub-array size in elements

The documentation for this class was generated from the following file:

- [dg_model_parameters.h](#)

6.8 DG::ModelParamsSection Struct Reference

Model parameters section descriptor.

```
#include <dg_model_parameters.h>
```

Public Attributes

- bool [is_scalar](#)
scalar vs vector flag: vector sections contain more than one element
- const char * [label](#)
section name string

6.8.1 Detailed Description

Model parameters section descriptor.

The documentation for this struct was generated from the following file:


- [dg_model_parameters.h](#)

6.9 DG::ModelParamsWriteAccess Class Reference


[ModelParamsWriteAccess](#) is read/write accessor to model parameters.

```
#include <dg_model_parameters.h>
```

Inheritance diagram for DG::ModelParamsWriteAccess:



Collaboration diagram for DG::ModelParamsWriteAccess:



Public Member Functions

- `bool is_dirty () const`
Check if dirty: at least one of the parameters was changed.
- `ModelParamsWriteAccess & merge (const json &config)`
- `ModelParamsWriteAccess (json &config)`
- `void set_dirty (bool state)`

Protected Member Functions

- `template<typename T>`
[ModelParamsWriteAccess](#) & [paramSet](#) (const char *section, const char *key, const T &value, size_t idx)

Protected Attributes

- [json](#) & [m_cfg_rw](#)
non-owning reference to JSON array with model configuration
- `bool` [m_dirty](#)
'some parameter was changed' flag

6.9.1 Detailed Description

[ModelParamsWriteAccess](#) is read/write accessor to model parameters.

This class provides programmatic type-safe read and write access to model parameters defined in JSON model configuration. It keeps non-owning non-const reference to underlying JSON array. For each model parameter it provides both getter and setter methods. Getter method name matches the parameter name as it appears in JSON array. Setter method name is constructed from the parameter name by appending `_set` suffix. For example, `InputImgFmt_set()`

6.9.2 Constructor & Destructor Documentation

6.9.2.1 ModelParamsWriteAccess()

```
DG::ModelParamsWriteAccess::ModelParamsWriteAccess (
    json & config ) [inline]
```

Constructor. Attaches model parameter read-write accessor to JSON array

Parameters

<code>in</code>	<code>config</code>	- JSON array with model configuration
-----------------	---------------------	---------------------------------------

6.9.3 Member Function Documentation

6.9.3.1 merge()

```
ModelParamsWriteAccess& DG::ModelParamsWriteAccess::merge (
    const json & config ) [inline]
```

Merge given model configuration JSON array with self Note: only runtime parameters will be merged

Parameters

<i>in</i>	<i>config</i>	- model configuration JSON array to be merged
-----------	---------------	---

6.9.3.2 paramSet()

```
template<typename T >
ModelParamsWriteAccess& DG::ModelParamsWriteAccess::paramSet (
    const char * section,
    const char * key,
    const T & value,
    size_t idx ) [inline], [protected]
```

Set parameter to JSON array

Template Parameters

<i>T</i>	- parameter type
----------	------------------

Parameters

<i>in</i>	<i>section</i>	- top section name; if empty, key is placed in topmost level
<i>in</i>	<i>key</i>	- key of key-value pair
<i>in</i>	<i>value</i>	- value to set
<i>in</i>	<i>idx</i>	- array index inside section object

6.9.3.3 set_dirty()

```
void DG::ModelParamsWriteAccess::set_dirty (
    bool state ) [inline]
```

Set dirty flag

Parameters

<i>in</i>	<i>state</i>	- dirty flag value to set
-----------	--------------	---------------------------

The documentation for this class was generated from the following file:

- [dg_model_parameters.h](#)

6.10 DG::ModelQuery Struct Reference

Model query structure used to search models on AI server which match a set of provided model attributes.

```
#include <dg_model_api.h>
```

Public Types

- enum [TriState](#) { [Yes](#), [No](#), [Dont_care](#) }
Tri-state boolean enum.

Public Attributes

- std::string [device_type](#)
device type to use for inference (optional, can be empty if don't care)
- std::string [model_name](#)
any part of the model name (mandatory parameter)
- [TriState](#) [model_pruned](#) = [Dont_care](#)
model should be pruned/sparse (not dense)
- [TriState](#) [model_quantized](#) = [Dont_care](#)
model should be quantized
- std::string [runtime_agent](#)
runtime agent to use for inference (optional, can be empty if don't care)

6.10.1 Detailed Description

Model query structure used to search models on AI server which match a set of provided model attributes.

6.10.2 Member Enumeration Documentation

6.10.2.1 TriState

```
enum DG::ModelQuery::TriState
```

Tri-state boolean enum.

Enumerator

Yes	yes / true
No	no / false
Dont_care	not set / don't care

The documentation for this struct was generated from the following file:

- [dg_model_api.h](#)

6.11 DG::BasicTensor::quant_params_t Struct Reference

Quantization parameters (how to convert from integer/quantized data back to floating point)

```
#include <dg_tensor_structs.h>
```

Classes

- struct [scale_t](#)

Public Member Functions

- bool [isEqualData](#) (const [quant_params_t](#) &rhs, double maxRelDiff) const
- bool [isEqualStruct](#) (const [quant_params_t](#) &rhs) const
- int [quant_axis](#) () const
Get quantization axis (-1 means global quantization)
- const std::vector< [scale_t](#) > & [quant_params](#) () const
Get quantization parameters.
- [quant_params_t](#) ()
Default constructor.
- [quant_params_t](#) (const [scale_t](#) &global_qparam)
- [quant_params_t](#) (int axis, const std::vector< [scale_t](#) > &qparams)
- template<typename SCALE_T, typename ZERO_T >
[quant_params_t](#) (int axis, const std::vector< SCALE_T > &scales, const std::vector< ZERO_T > &zeros)
- template<typename T >
std::vector< T > [quant_scales](#) () const
Get array of scales.
- template<typename T >
std::vector< T > [quant_zeros](#) () const
Get array of zero offsets.

6.11.1 Detailed Description

Quantization parameters (how to convert from integer/quantized data back to floating point)

6.11.2 Constructor & Destructor Documentation

6.11.2.1 quant_params_t() [1/3]

```
DG::BasicTensor::quant_params_t::quant_params_t (  
    const scale\_t & global_qparam ) [inline]
```

Constructor for global quantization

Parameters

in	<i>global_qparam</i>	- global quantization parameters
----	----------------------	----------------------------------

6.11.2.2 quant_params_t() [2/3]

```
DG::BasicTensor::quant_params_t::quant_params_t (
    int axis,
    const std::vector< scale_t > & qparams ) [inline]
```

Constructor for per-axis quantization

Parameters

in	<i>axis</i>	- quantization axis
in	<i>qparams</i>	- vector of per-axis quantization parameters

6.11.2.3 quant_params_t() [3/3]

```
template<typename SCALE_T , typename ZERO_T >
DG::BasicTensor::quant_params_t::quant_params_t (
    int axis,
    const std::vector< SCALE_T > & scales,
    const std::vector< ZERO_T > & zeros ) [inline]
```

Alternative constructor for per-axis quantization

Parameters

in	<i>axis</i>	- quantization axis
in	<i>scales</i>	- vector of per-axis scales
in	<i>zeros</i>	- vector of per-axis zero offsets

6.11.3 Member Function Documentation**6.11.3.1 isEqualData()**

```
bool DG::BasicTensor::quant_params_t::isEqualData (
    const quant_params_t & rhs,
    double maxRelDiff ) const [inline]
```

Check, if this object and given object have equal data contents

Parameters

in	<i>rhs</i>	- object to compare with
in	<i>maxRelDiff</i>	- maximum relative difference for floating point comparison

6.11.3.2 isEqualStruct()

```
bool DG::BasicTensor::quant_params_t::isEqualStruct (
    const quant_params_t & rhs ) const [inline]
```

Check, if this object and given object have equal structure

Parameters

in	<i>rhs</i>	- object to compare with
----	------------	--------------------------

The documentation for this struct was generated from the following file:

- [dg_tensor_structs.h](#)

6.12 DG::BasicTensor::quant_params_t::scale_t Struct Reference

```
#include <dg_tensor_structs.h>
```

Public Attributes

- double [m_scale](#)
scale factor
- int64_t [m_zero](#)
zero offset

6.12.1 Detailed Description

Single quantization parameter: scale and zero offset Scaling formula is $\text{real_value} = \text{m_scale} * (\text{int_value} - \text{m_zero})$

The documentation for this struct was generated from the following file:

- [dg_tensor_structs.h](#)

6.13 DG::ServerAddress Struct Reference

```
#include <dg_client_structs.h>
```

Public Member Functions

- const bool [is_valid](#) () const
- [operator std::string](#) () const
Convert server address to string.
- [ServerAddress](#) (const std::string &[ip](#), int [port](#), [ServerType](#) [server_type](#))

Static Public Member Functions

- static [ServerAddress](#) [fromHostname](#) (const std::string &hostname)

Public Attributes

- std::string [ip](#)
server domain name or IP address string
- int [port](#) = [DEFAULT_PORT](#)
server TCP port number
- [ServerType](#) [server_type](#) = [ServerType::Unknown](#)
server protocol type

6.13.1 Detailed Description

[ServerAddress](#) is the server address structure. It keeps AI server TCP/IP address, port, and server protocol type.

6.13.2 Constructor & Destructor Documentation

6.13.2.1 [ServerAddress\(\)](#)

```
DG::ServerAddress::ServerAddress (
    const std::string & ip,
    int port,
    ServerType server_type ) [inline]
```

Constructor

Parameters

in	<i>ip</i>	- server domain name or IP address
in	<i>port</i>	- server TCP port
in	<i>server_type</i>	- server protocol type

6.13.3 Member Function Documentation

6.13.3.1 fromHostname()

```
static ServerAddress DG::ServerAddress::fromHostname (
    const std::string & hostname ) [inline], [static]
```

Construct server address from hostname

Parameters

in	<i>hostname</i>	- server domain name or IP address with optional port suffix and protocol prefix (http:// or asio://)
----	-----------------	--

Returns

server address object

6.13.3.2 is_valid()

```
const bool DG::ServerAddress::is_valid ( ) const [inline]
```

Check server address validity

Returns

true if server address is valid

The documentation for this struct was generated from the following file:

- [dg_client_structs.h](#)


Chapter 7

File Documentation


7.1 dg_client_structs.h File Reference

DG client API data types.

```
#include <stdlib.h>
#include <string>
#include <vector>
#include "Utilities/dg_model_parameters.h"
Include dependency graph for dg_client_structs.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- struct [DG::ModelInfo](#)
ModelInfo is the model identification structure. It keeps AI model key attributes.
- struct [DG::ServerAddress](#)

Namespaces

- [DG](#)
DG is DeGirum namespace which groups all DeGirum software declarations.

Typedefs

- typedef struct [DG::ModelInfo](#) [DG::ModelInfo](#)
ModelInfo is the model identification structure. It keeps AI model key attributes.

Enumerations

- enum [DG::ServerType](#) { [DG::ServerType::Unknown](#), [DG::ServerType::ASIO](#), [DG::ServerType::HTTP](#) }
Server protocol type.

Functions

- static std::string [DG::messagePrepare](#) (const json &in)
- static json [DG::messagePrepareJson](#) (const json &in)

Variables

- const int [DG::CURRENT_PROTOCOL_VERSION](#) = 4
current client-server protocol version
- const int [DG::DEFAULT_CONNECTION_TIMEOUT_MS](#) = 10000
Default connection timeout, ms.
- const int [DG::DEFAULT_FRAME_QUEUE_DEPTH](#) = 8
Default frame queue depth.
- const int [DG::DEFAULT_INFERENCE_TIMEOUT_MS](#) = 180000
Default inference timeout, ms.
- const int [DG::DEFAULT_PORT](#) = 8778
Default TCP port of AI server.
- const int [DG::MIN_COMPATIBLE_PROTOCOL_VERSION](#) = 4
minimum compatible client-server protocol version
- constexpr const char * [DG::PROTOCOL_VERSION_TAG](#) = "VERSION"
client-server protocol version tag

7.1.1 Detailed Description

DG client API data types.

This file contains declarations of various data types of DG client API:


- device types
- runtime agent types
- model parameters, etc.

7.2 dg_json_helpers.h File Reference


DG Core JSON helper classes and functions.

```
#include "Utilities/DGErrorHandling.h"
#include "Utilities/dg_tensor_structs.h"
#include "Utilities/dg_raii_helpers.h"
#include "Utilities/json.hpp"
#include <functional>
#include <algorithm>
```

Include dependency graph for dg_json_helpers.h:



This graph shows which files directly or indirectly include this file:



Classes

- class [DG::JsonHelper](#)

JSON helper class: contains assorted static methods to operate with JSON configurations.

Namespaces

- [DG](#)

DG is DeGirum namespace which groups all DeGirum software declarations.

Macros

- `#define DG\_JSON\_PARSE(json_cfg) DG::JsonHelper::parse( json_cfg, __FILE__, TOSTRING(__LINE__), FUNCTION_NAME )`

Parse given JSON string with proper exception handling.

Typedefs

- using [DG::json](#) = nlohmann::json

JSON library.

Functions

- `template<typename T >`
`T DG::jsonGetMandatoryValue (const json &json_params, const std::string §ion, int index, const std::string &key)`
- `template<typename T >`
`T DG::jsonGetOptionalValue (const json &json_params, const std::string §ion, int index, const std::string &key, const T &default_value)`
- `static bool DG::jsonKeyExist (const json &json_params, const std::string §ion, int index, const std::string &key)`
- `template<typename T >`
`void DG::jsonSetOptionalValue (json &json_params, const std::string §ion, int index, const std::string &key, const T &value)`

7.2.1 Detailed Description


[DG](#) Core JSON helper classes and functions.

This file contains declaration of various helper classes and helper functions enhancing JSON handling

7.3 dg_model_api.h File Reference

DG Client API for model inference.

```
#include <stdint>
#include <functional>
#include <map>
#include <memory>
#include <set>
#include <string>
#include <vector>
#include "Utilities/dg_client_structs.h"
#include "Utilities/dg_model_parameters.h"
Include dependency graph for dg_model_api.h:
```



Classes

- class [DG::AIModel](#)
AIModel is DeGirum AI client API class for simple non-pipelined sequential inference.
- class [DG::AIModelAsync](#)
AIModelAsync is DeGirum AI client API class for efficient pipelined asynchronous inference.
- struct [DG::ModelQuery](#)
Model query structure used to search models on AI server which match a set of provided model attributes.

Namespaces

- [DG](#)
DG is DeGirum namespace which groups all DeGirum software declarations.

Enumerations

- enum [DG::DetectionStatus](#) { [DG::DetectionStatus::OK](#), [DG::DetectionStatus::ProtocolMismatch](#), [DG::DetectionStatus::HostAlive](#) }
Enumerator showing server status as detected by the below function.

Functions

- `std::vector< std::tuple< std::string, DG::DetectionStatus > > DG::detectHostnameServers` (const std::string &prefix, const int range_start, const int range_end, const int numeral_width=3)
- `std::vector< std::tuple< std::string, DetectionStatus > > DG::detectSubnetServers` (const std::string &root_ip, const std::string &subnet_mask)
- `json DG::devCtrl` (const std::string &server, const json &req)
- `std::string DG::errorCheck` (const json &json_response)
- `json DG::labelDictionary` (const std::string &server, const std::string &model_name)
- `ModelInfo DG::modelFind` (const std::string &server, const ModelQuery &query)
- `void DG::modelzooListGet` (const std::string &server, std::vector< ModelInfo > &modelzoo_list)
- `json DG::modelZooManage` (const std::string &server, const json &req)
- `bool DG::serverPing` (const std::string &server)
- `void DG::shutdown` (const std::string &server)
- `json DG::systemInfo` (const std::string &server)
- `json DG::traceManage` (const std::string &server, const json &req)
- `std::string DG::versionGet` ()

7.3.1 Detailed Description


[DG](#) Client API for model inference.

This file contains declaration of [DG](#) Client API functions and classes for model inference


7.4 dg_model_parameters.h File Reference

[DG](#) centralized handling of JSON model parameters.

```
#include <string>
#include <type_traits>
#include "dg_json_helpers.h"
#include "dg_model_parameters.inc"
Include dependency graph for dg_model_parameters.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class [DG::ModelParams< Base, >](#)
ModelParams is model parameters collection with user-defined access rights.
- class [DG::ModelParamsReadAccess](#)
ModelParamsReadAccess is read-only accessor to model parameters.
- struct [DG::ModelParamsSection](#)
Model parameters section descriptor.
- class [DG::ModelParamsWriteAccess](#)
ModelParamsWriteAccess is read/write accessor to model parameters.

Namespaces

- [DG](#)
DG is DeGirum namespace which groups all DeGirum software declarations.

Typedefs

- using [DG::ModelParamsWriter](#) = [ModelParams< ModelParamsWriteAccess >](#)
ModelParamsWriter is ModelParams template instantiation with write access.
- using [DG::vec3d_int](#) = [std::vector< std::vector< std::vector< int > > >](#)
3D vector of integers
- using [DG::vec_dbl](#) = [std::vector< double >](#)
vector of doubles
- using [DG::vecflt](#) = [std::vector< float >](#)
vector of floats
- using [DG::vec_int](#) = [std::vector< int >](#)
vector of integers
- using [DG::vec_str](#) = [std::vector< std::string >](#)
vector of strings
- using [DG::vec_sz](#) = [std::vector< size_t >](#)
vector of size_t

Variables

- const int [DG::MODEL_PARAMS_CURRENT_VERSION](#) = 12
- const int [DG::MODEL_PARAMS_MIN_COMPATIBLE_VERSION](#) = 1
- constexpr ModelParamsSection [DG::SECT_DEVICE](#) = { "DEVICE", true }
device parameters section
- constexpr ModelParamsSection [DG::SECT_INTERNAL](#) = { "INTERNAL", true }
internal parameters section
- constexpr ModelParamsSection [DG::SECT_MODEL_PARAMETERS](#) = { "MODEL_PARAMETERS", true }
model parameters section
- constexpr ModelParamsSection [DG::SECT_POST_PROCESS](#) = { "POST_PROCESS", true }
post-processing parameters section
- constexpr ModelParamsSection [DG::SECT_PRE_PROCESS](#) = { "PRE_PROCESS", false }
pre-processing parameters section
- constexpr ModelParamsSection [DG::SECT_TOP](#) = { "", true }
JSON file top-level sections.

7.4.1 Detailed Description

[DG](#) centralized handling of JSON model parameters.


This file implements classes for centralized handling of JSON model parameters. Each AI model in [DG](#) framework is accompanied with JSON configuration file, which defines all model configuration parameters. Class list: ModelParamsReadAccess - read-only model parameters accessor ModelParamsWriteAccess - read/write model parameters accessor ModelParams - model parameters collection with access type defined by the template parameter

7.5 dg_tensor_structs.h File Reference


[DG](#) tensor container classes.

```
#include <algorithm>
#include <numeric>
#include <string>
#include <typeinfo>
#include <variant>
#include <vector>
#include "Utilities/DGErrorHandling.h"
#include "Utilities/TypeList.h"
#include "Utilities/dg_math_utilities.h"
#include "Utilities/dg_raii_helpers.h"
#include "Utilities/nameof.hpp"
```

Include dependency graph for `dg_tensor_structs.h`:



This graph shows which files directly or indirectly include this file:



Classes

- class [DG::BasicTensor](#)
- struct [DG::BasicTensor::quant_params_t](#)
Quantization parameters (how to convert from integer/quantized data back to floating point)
- struct [DG::BasicTensor::quant_params_t::scale_t](#)

Namespaces

- [DG](#)
DG is DeGirum namespace which groups all DeGirum software declarations.

Typedefs

- using [DG::BasicTensorVector](#) = `std::vector< BasicTensor >`
Collection of tensors.

Functions

- `std::ostream & DG::operator<< (std::ostream &os, const BasicTensor &t)`
Print tensor to stream.

7.5.1 Detailed Description

[DG](#) tensor container classes.

This file contains declaration of tensor container classes

Index

- ~AIModelAsync
 - DG::AIModelAsync, [38](#)
- adopt
 - DG::BasicTensor, [44](#)
- AIModel
 - DG::AIModel, [35](#)
- AIModelAsync
 - DG::AIModelAsync, [38](#)
- alloc
 - DG::BasicTensor, [45](#)
- ASIO
 - DG, [26](#)
- BasicTensor
 - DG::BasicTensor, [43](#)
- callback_t
 - DG::AIModelAsync, [37](#)
- convert
 - DG::BasicTensor, [46](#)
- copy
 - DG::BasicTensor, [47](#)
- data
 - DG::BasicTensor, [47](#)
- dequantize
 - DG::BasicTensor, [47](#)
- detectHostnameServers
 - DG, [26](#)
- DetectionStatus
 - DG, [25](#)
- detectSubnetServers
 - DG, [27](#)
- devCtrl
 - DG, [27](#)
- DG, [23](#)
 - ASIO, [26](#)
 - detectHostnameServers, [26](#)
 - DetectionStatus, [25](#)
 - detectSubnetServers, [27](#)
 - devCtrl, [27](#)
 - errorCheck, [28](#)
 - HostAlive, [26](#)
 - HTTP, [26](#)
 - jsonGetMandatoryValue, [28](#)
 - jsonGetOptionalValue, [28](#)
 - jsonKeyExist, [29](#)
 - jsonSetOptionalValue, [29](#)
 - labelDictionary, [30](#)
 - messagePrepare, [30](#)
 - messagePrepareJson, [30](#)
 - MODEL_PARAMS_CURRENT_VERSION, [34](#)
 - MODEL_PARAMS_MIN_COMPATIBLE_VERSION, [34](#)
 - modelFind, [31](#)
 - modelzooListGet, [31](#)
 - modelZooManage, [32](#)
 - OK, [26](#)
 - ProtocolMismatch, [26](#)
 - SECT_TOP, [34](#)
 - serverPing, [32](#)
 - ServerType, [26](#)
 - shutdown, [32](#)
 - systemInfo, [33](#)
 - traceManage, [33](#)
 - Unknown, [26](#)
 - versionGet, [33](#)
- DG::AIModel, [35](#)
 - AIModel, [35](#)
 - predict, [36](#)
- DG::AIModelAsync, [37](#)
 - ~AIModelAsync, [38](#)
 - AIModelAsync, [38](#)
 - callback_t, [37](#)
 - lastError, [39](#)
 - predict, [39](#)
 - setCallback, [39](#)
 - waitCompletion, [40](#)
- DG::BasicTensor, [40](#)
 - adopt, [44](#)
 - alloc, [45](#)
 - BasicTensor, [43](#)
 - convert, [46](#)
 - copy, [47](#)
 - data, [47](#)
 - dequantize, [47](#)
 - isEqualData, [47](#), [48](#)
 - isEqualDataShape, [48](#)
 - isEqualStruct, [48](#)
 - linearSizeCalc, [48](#)
 - quantize, [49](#)
 - quantParamsSet, [49](#)
 - reinterpretShapeForYolo, [49](#)
 - reinterpretShapeScaled, [49](#)
 - reshapeCombineDims, [49](#)
 - reshapeTo, [50](#)
 - shapeStringGet, [50](#)
- DG::BasicTensor::quant_params_t, [69](#)

- isEqualData, [70](#)
- isEqualStruct, [71](#)
- quant_params_t, [69](#), [70](#)
- DG::BasicTensor::quant_params_t::scale_t, [71](#)
- DG::JsonHelper, [51](#)
 - errorCheck, [51](#)
 - jsonDeserialize, [52](#)
 - jsonDeserializeStr, [52](#)
 - jsonSerialize, [52](#)
 - jsonSerializeStr, [53](#)
 - parse, [53](#)
 - parse_ignore_errors, [54](#)
 - tensorDeserialize, [54](#)
 - tensorSerialize, [54](#)
- DG::ModelInfo, [55](#)
- DG::ModelParams< Base, >, [56](#)
 - ModelParams, [57](#), [58](#)
- DG::ModelParamsReadAccess, [59](#)
 - fullJsonGet, [60](#)
 - jsonGet, [60](#)
 - ModelParamsReadAccess, [60](#)
 - modelShapeGet, [61](#)
 - operator const json &, [61](#)
 - operator std::string, [61](#)
 - paramExist, [61](#)
 - paramGet, [62](#)
 - sectionGet, [63](#)
 - sectionHashGet, [63](#)
 - sectionSizeGet, [63](#)
- DG::ModelParamsSection, [64](#)
- DG::ModelParamsWriteAccess, [65](#)
 - merge, [66](#)
 - ModelParamsWriteAccess, [66](#)
 - paramSet, [67](#)
 - set_dirty, [67](#)
- DG::ModelQuery, [67](#)
 - Dont_care, [68](#)
 - No, [68](#)
 - TriState, [68](#)
 - Yes, [68](#)
- DG::ServerAddress, [71](#)
 - fromHostname, [73](#)
 - is_valid, [73](#)
 - ServerAddress, [72](#)
- dg_client_structs.h, [75](#)
- dg_json_helpers.h, [77](#)
- dg_model_api.h, [79](#)
- dg_model_parameters.h, [80](#)
- dg_tensor_structs.h, [82](#)
- Dont_care
 - DG::ModelQuery, [68](#)
- errorCheck
 - DG, [28](#)
 - DG::JsonHelper, [51](#)
- fromHostname
 - DG::ServerAddress, [73](#)
- fullJsonGet
 - DG::ModelParamsReadAccess, [60](#)
- HostAlive
 - DG, [26](#)
- HTTP
 - DG, [26](#)
- is_valid
 - DG::ServerAddress, [73](#)
- isEqualData
 - DG::BasicTensor, [47](#), [48](#)
 - DG::BasicTensor::quant_params_t, [70](#)
- isEqualDataShape
 - DG::BasicTensor, [48](#)
- isEqualStruct
 - DG::BasicTensor, [48](#)
 - DG::BasicTensor::quant_params_t, [71](#)
- jsonDeserialize
 - DG::JsonHelper, [52](#)
- jsonDeserializeStr
 - DG::JsonHelper, [52](#)
- jsonGet
 - DG::ModelParamsReadAccess, [60](#)
- jsonGetMandatoryValue
 - DG, [28](#)
- jsonGetOptionalValue
 - DG, [28](#)
- jsonKeyExist
 - DG, [29](#)
- jsonSerialize
 - DG::JsonHelper, [52](#)
- jsonSerializeStr
 - DG::JsonHelper, [53](#)
- jsonSetOptionalValue
 - DG, [29](#)
- labelDictionary
 - DG, [30](#)
- lastError
 - DG::AIModelAsync, [39](#)
- linearSizeCalc
 - DG::BasicTensor, [48](#)
- merge
 - DG::ModelParamsWriteAccess, [66](#)
- messagePrepare
 - DG, [30](#)
- messagePrepareJson
 - DG, [30](#)
- MODEL_PARAMS_CURRENT_VERSION
 - DG, [34](#)
- MODEL_PARAMS_MIN_COMPATIBLE_VERSION
 - DG, [34](#)
- modelFind
 - DG, [31](#)
- ModelParams
 - DG::ModelParams< Base, >, [57](#), [58](#)
- ModelParamsReadAccess

- DG::ModelParamsReadAccess, 60
- ModelParamsWriteAccess
 - DG::ModelParamsWriteAccess, 66
- modelShapeGet
 - DG::ModelParamsReadAccess, 61
- modelzooListGet
 - DG, 31
- modelZooManage
 - DG, 32
- No
 - DG::ModelQuery, 68
- OK
 - DG, 26
- operator const json &
 - DG::ModelParamsReadAccess, 61
- operator std::string
 - DG::ModelParamsReadAccess, 61
- paramExist
 - DG::ModelParamsReadAccess, 61
- paramGet
 - DG::ModelParamsReadAccess, 62
- paramSet
 - DG::ModelParamsWriteAccess, 67
- parse
 - DG::JsonHelper, 53
- parse_ignore_errors
 - DG::JsonHelper, 54
- predict
 - DG::AIModel, 36
 - DG::AIModelAsync, 39
- ProtocolMismatch
 - DG, 26
- quant_params_t
 - DG::BasicTensor::quant_params_t, 69, 70
- quantize
 - DG::BasicTensor, 49
- quantParamsSet
 - DG::BasicTensor, 49
- reinterpretShapeForYolo
 - DG::BasicTensor, 49
- reinterpretShapeScaled
 - DG::BasicTensor, 49
- reshapeCombineDims
 - DG::BasicTensor, 49
- reshapeTo
 - DG::BasicTensor, 50
- SECT_TOP
 - DG, 34
- sectionGet
 - DG::ModelParamsReadAccess, 63
- sectionHashGet
 - DG::ModelParamsReadAccess, 63
- sectionSizeGet
 - DG::ModelParamsReadAccess, 63
- ServerAddress
 - DG::ServerAddress, 72
- serverPing
 - DG, 32
- ServerType
 - DG, 26
- set_dirty
 - DG::ModelParamsWriteAccess, 67
- setCallback
 - DG::AIModelAsync, 39
- shapeStringGet
 - DG::BasicTensor, 50
- shutdown
 - DG, 32
- systemInfo
 - DG, 33
- tensorDeserialize
 - DG::JsonHelper, 54
- tensorSerialize
 - DG::JsonHelper, 54
- traceManage
 - DG, 33
- TriState
 - DG::ModelQuery, 68
- Unknown
 - DG, 26
- versionGet
 - DG, 33
- waitCompletion
 - DG::AIModelAsync, 40
- Yes
 - DG::ModelQuery, 68